Recurrent Neural Networks (RNN)

Jimeng Sun

ILLINOIS

Recurrent Neural Networks

Agenda

- RNN Basics
- Learning RNN with Backpropagation Through Time (BPTT)
- Long-Short Term Memory Networks (LSTM)
- Gated Recurrent Unit (GRU)
- Bidirectional RNN
- Sequence-to-Sequence RNN
- Healthcare Applications

Agenda

- RNN Basics
- Learning RNN with Backpropagation Through Time (BPTT)
- Long-Short Term Memory Networks (LSTM)
- Gated Recurrent Unit (GRU)
- Bidirectional RNN
- Sequence-to-Sequence RNN
- Healthcare Applications

ILLINOIS

Basic Concepts of RNN

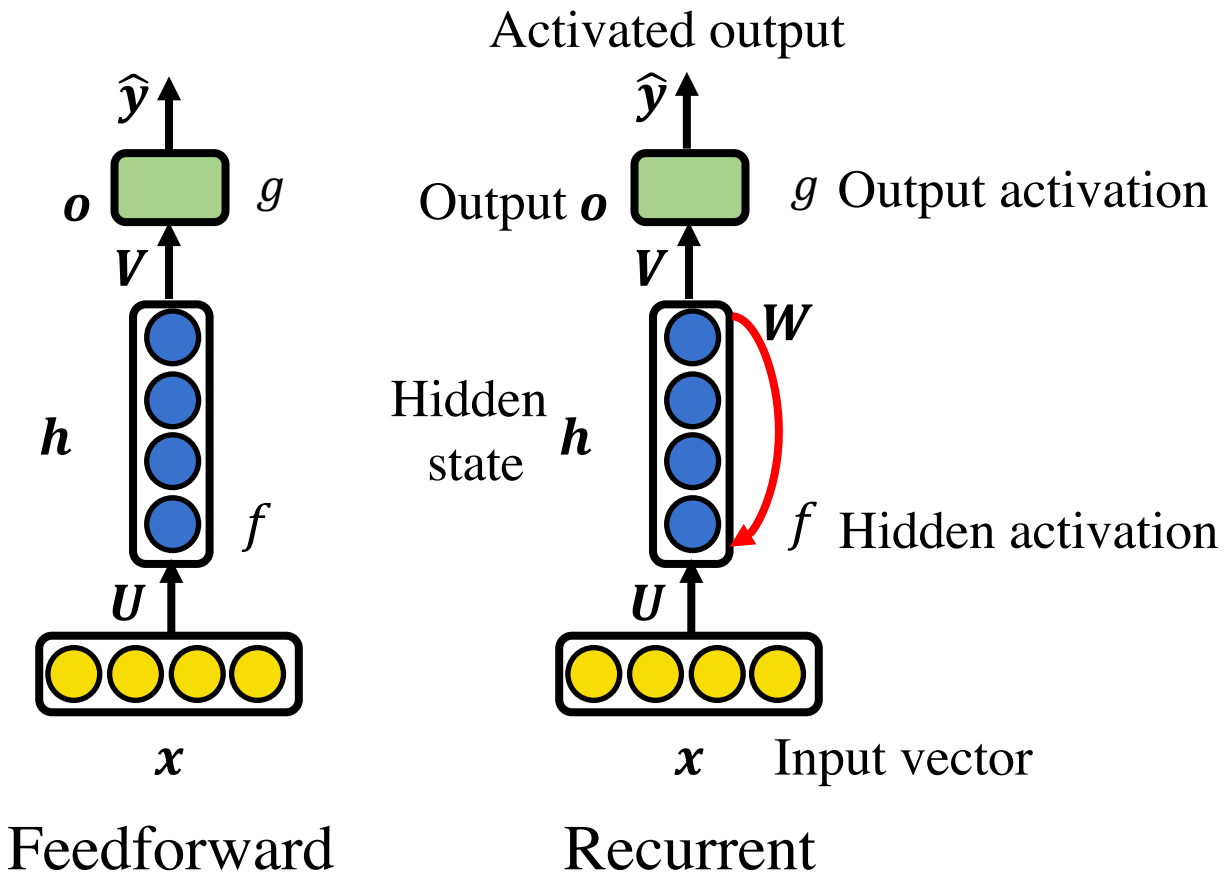

Basic RNN Structure

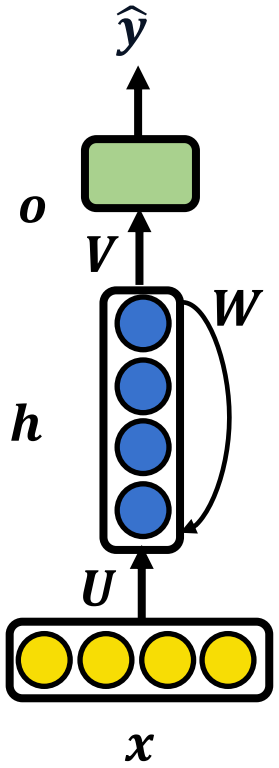

unroll

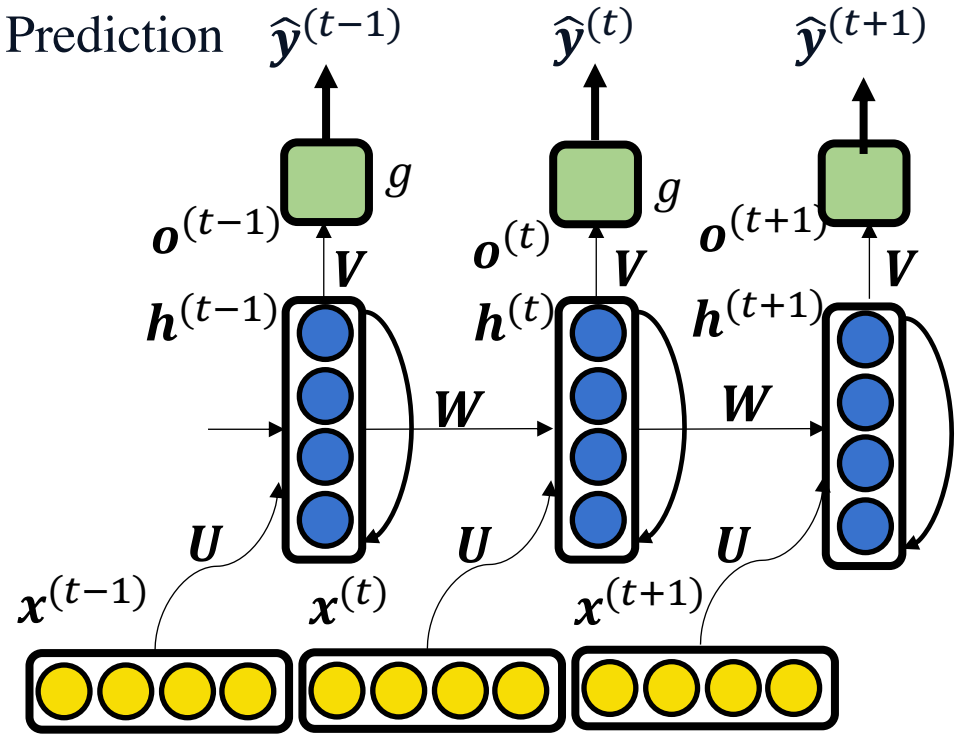

$$\boldsymbol{h}^{(t)} = f(\boldsymbol{U}\boldsymbol{x}^{(t)} + \boldsymbol{W}\boldsymbol{h}^{(t-1)} + \boldsymbol{b}_1)$$
$$\boldsymbol{o}^{(t)} = \boldsymbol{V}\boldsymbol{h}^{(t)} + \boldsymbol{b}_2$$
$$\hat{\boldsymbol{y}}^{(t)} = g(\boldsymbol{o}^{(t)})$$

ILLINOIS

Basic RNN Structure

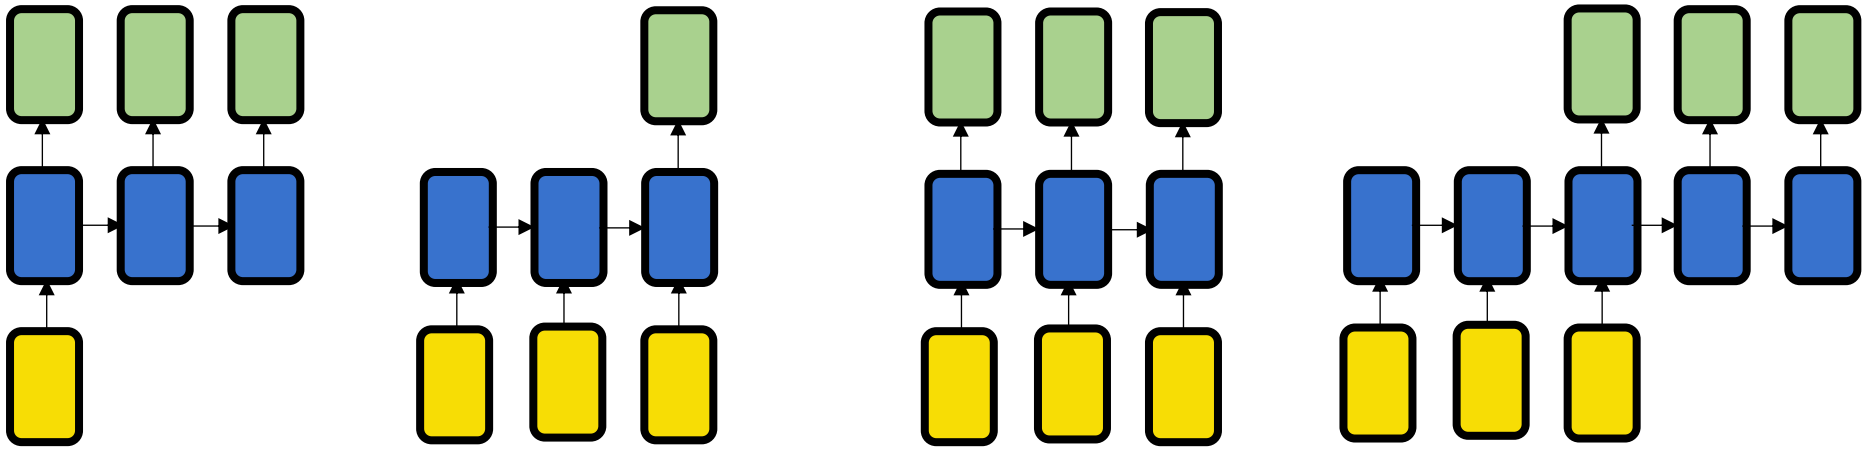

One-to-Many (e.g., image to text)

Many-to-One (e.g., sequence classification)

Many-to-Many (e.g., sequential prediction)

Many-to-Many (e.g., seq2seq)

The figure is inspired by page 12 on http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture10.pdf

Agenda

- RNN Basics
- Learning RNN with Backpropagation Through Time (BPTT)
- Long-Short Term Memory Networks (LSTM)
- Gated Recurrent Unit (GRU)
- Bidirectional RNN
- Sequence-to-Sequence RNN
- Healthcare Applications

ILLINOIS

Forward Computation

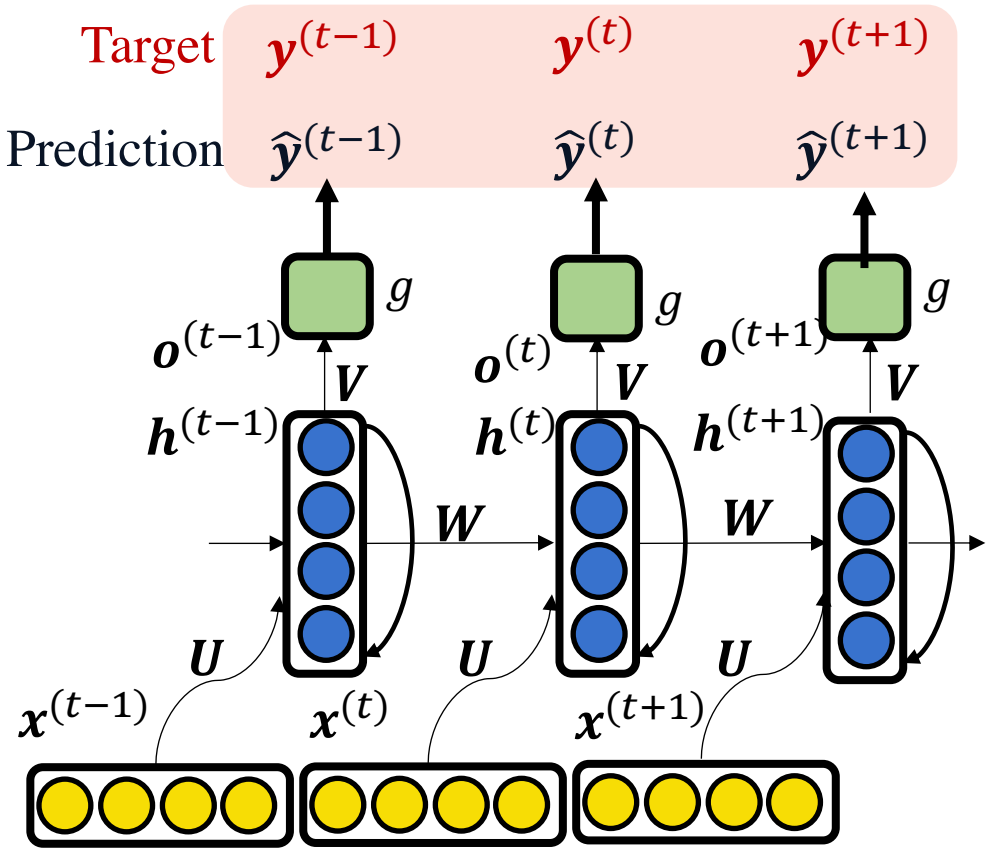

$$z^{(t)} = Ux^{(t)} + Wh^{(t-1)} + b_1$$
$$h^{(t)} = f(z^{(t)})$$
$$o^{(t)} = Vh^{(t)} + b_2$$
$$\hat{y}^{(t)} = g(o^{(t)})$$
$$L = \sum_t L^{(t)} = -\sum_t \log p(y^{(t)} | \{x^{(t)}, \cdots, x^{(T)}\})$$

e. g.

- f is tanh
- g is softmax, which produces a normalized probability over output classes
- $L^{(t)}$ is negative log-likelihood loss

e.g., in binary classification

$$L = -\sum_t y^{(t)} \log \hat{y}^{(t)} + (1 - y^{(t)}) \log(1 - \hat{y}^{(t)})$$

ILLINOIS

Backpropagation through time (BPTT): $\nabla_{\boldsymbol{o}^{(t)}} L$

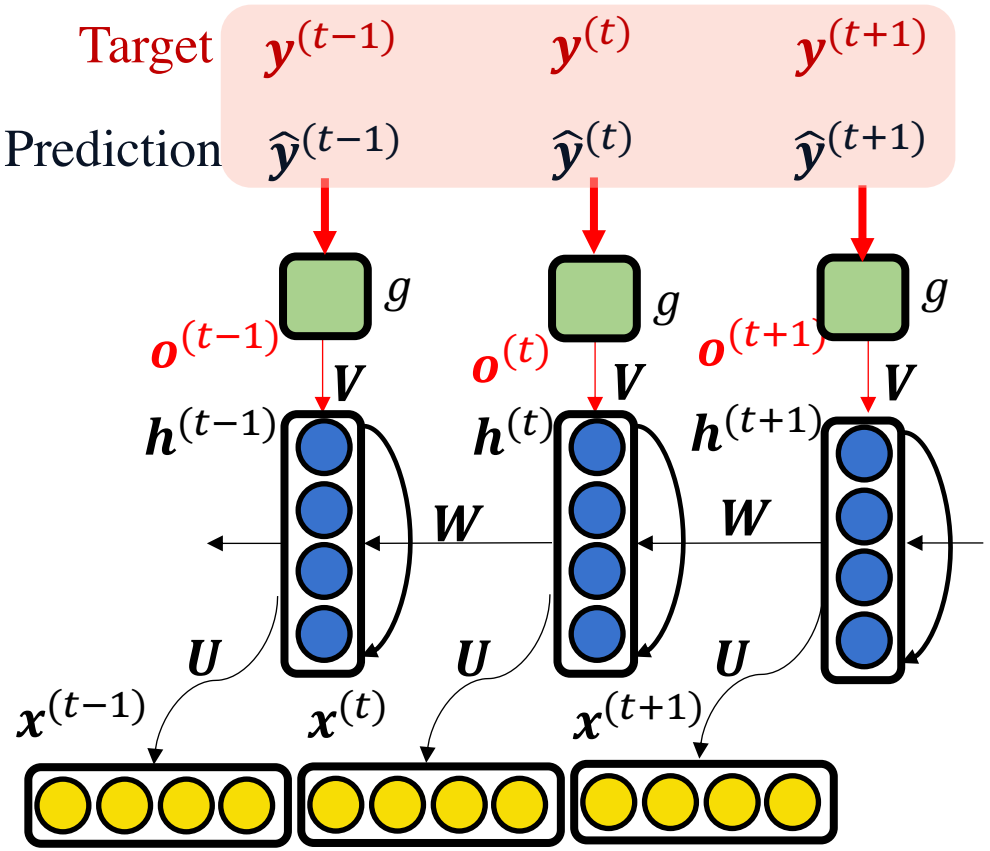

$$\left(\nabla_{\boldsymbol{o}^{(t)}} L\right)_i = \frac{\partial L}{\partial \boldsymbol{o}_i^{(t)}} = \frac{\partial L}{\partial L^{(t)}} \frac{\partial L^{(t)}}{\partial \boldsymbol{o}_i^{(t)}}$$

ILLINOIS

Backpropagation through time (BPTT): $\nabla_{h^{(T)}} L$

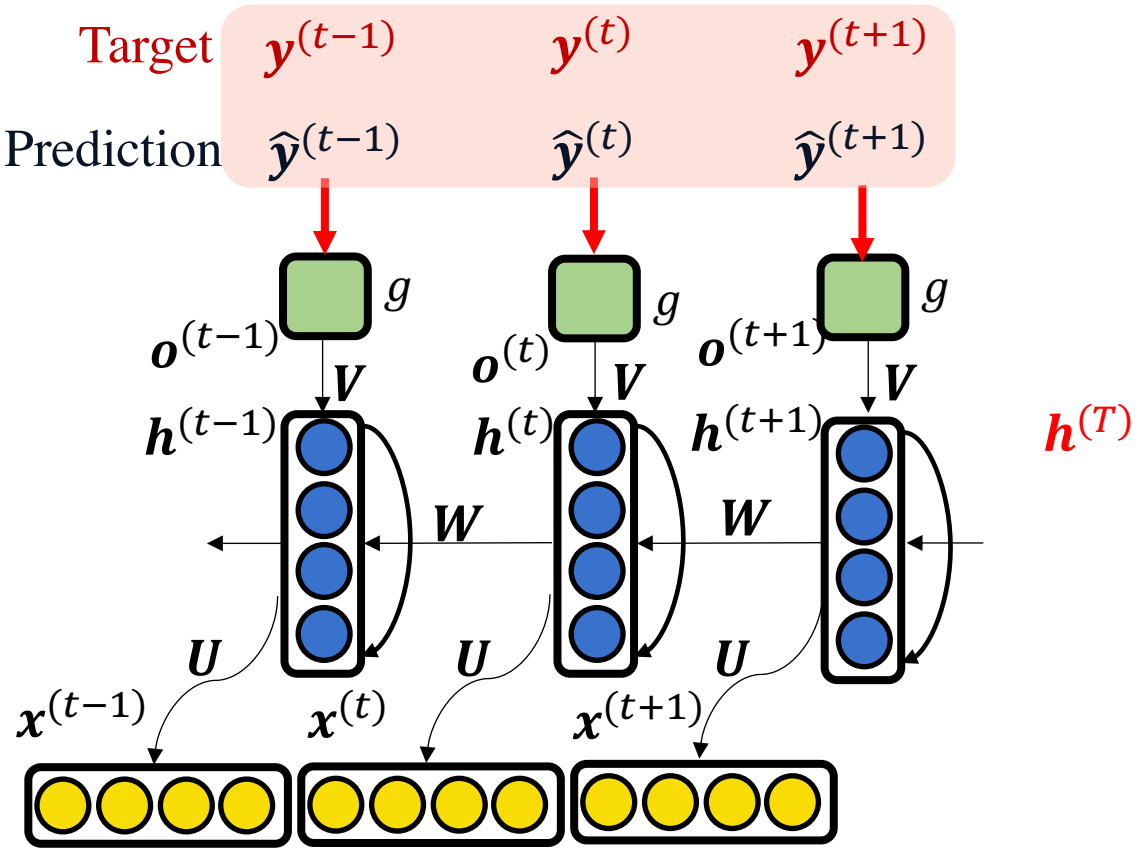

Last time stamp

$$\nabla_{h^{(T)}} L = V^T \nabla_{o^{(T)}} L$$

ILLINOIS

Backpropagation through time (BPTT): $\nabla_{\boldsymbol{h}^{(t)}} L$

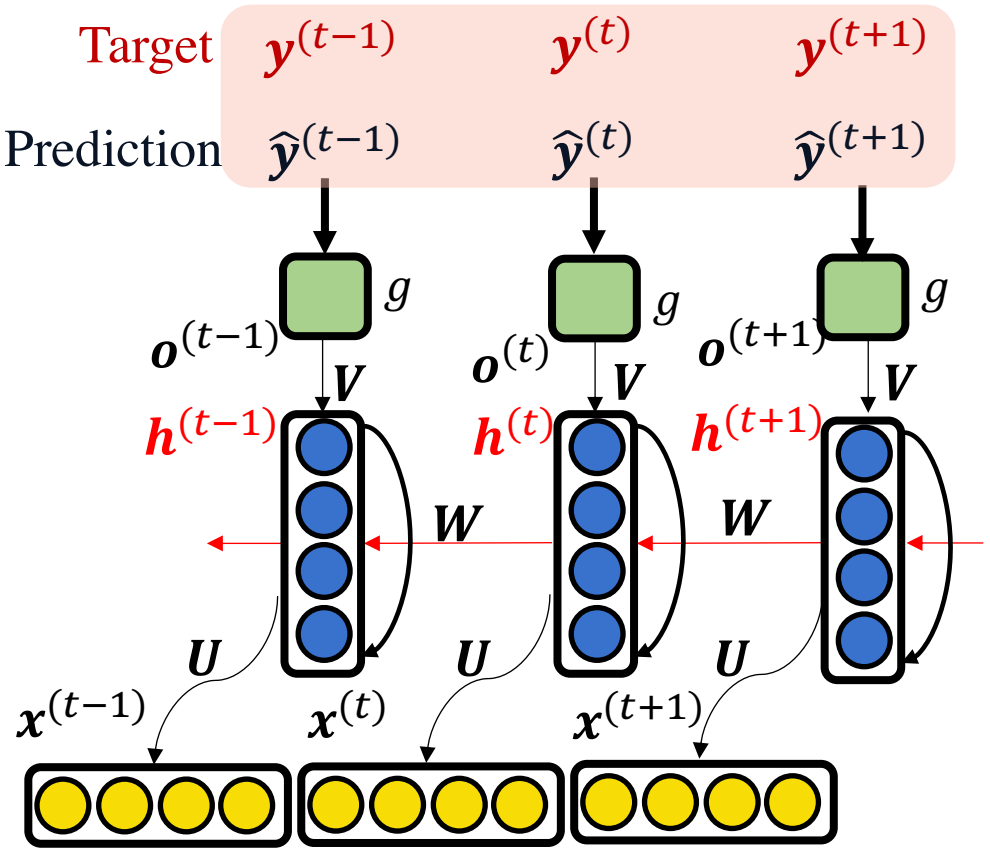

Last time stamp

$$\nabla_{\boldsymbol{h}^{(T)}} L = V^T \nabla_{\boldsymbol{o}^{(T)}} L$$

Other time stamp

$$\nabla_{\boldsymbol{h}^{(t)}} L = \left(\frac{\partial \boldsymbol{h}^{(t+1)}}{\partial \boldsymbol{h}^{(t)}}\right)^{\boldsymbol{T}} \left(\nabla_{\boldsymbol{h}^{(t+1)}} L\right) + \left(\frac{\partial \boldsymbol{o}^{(t)}}{\partial \boldsymbol{h}^{(t)}}\right)^{\boldsymbol{T}} \left(\nabla_{\boldsymbol{o}^{(t)}} L\right)$$

ILLINOIS

Backpropagation through time (BPTT)

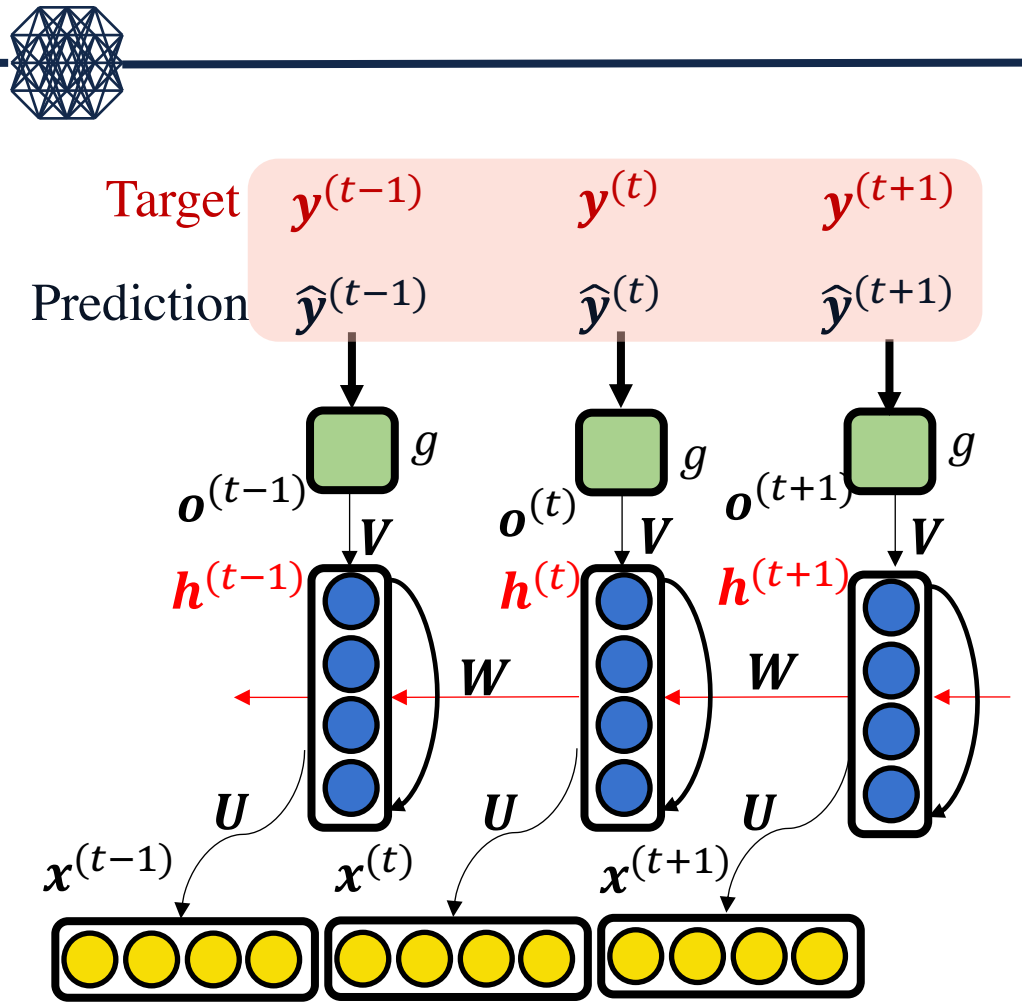

$$\nabla_{\boldsymbol{V}}\mathcal{L} = \sum_t \sum_k \left(\frac{\partial \mathcal{L}}{\partial o_k^{(t)}}\right) \nabla_{\boldsymbol{V}} o_k^{(t)} = \sum_t (\nabla_{\boldsymbol{o}^{(t)}}\mathcal{L}) \boldsymbol{h}^{(t)^T}$$

$$\nabla_{\boldsymbol{W}}\mathcal{L} = \sum_t \sum_j \left(\frac{\partial \mathcal{L}}{\partial h_j^{(t)}}\right) \nabla_{\boldsymbol{W}} h_j^{(t)} = \sum_t diag(1-(\boldsymbol{h}^{(t)})^2)(\nabla_{\boldsymbol{h}^{(t)}}\mathcal{L}) \boldsymbol{h}^{(t-1)^T}$$

$$\nabla_{\boldsymbol{U}}\mathcal{L} = \sum_t \sum_j \left(\frac{\partial \mathcal{L}}{\partial h_j^{(t)}}\right) \nabla_{\boldsymbol{U}} h_j^{(t)} = \sum_t diag(1-(\boldsymbol{h}^{(t)})^2)(\nabla_{\boldsymbol{h}^{(t)}}\mathcal{L}) \boldsymbol{x}^{(t)^T}$$

$$\nabla_{\boldsymbol{b}_1}\mathcal{L} = \sum_t \left(\frac{\partial \boldsymbol{h}^{(t)}}{\partial \boldsymbol{b}^{(t)}}\right)^T \nabla_{\boldsymbol{h}^{(t)}}\mathcal{L} = \sum_t diag(1-(\boldsymbol{h}^{(t)})^2) \nabla_{\boldsymbol{h}^{(t)}}\mathcal{L}$$

$$\nabla_{\boldsymbol{b}_2}\mathcal{L} = \sum_t \left(\frac{\partial \boldsymbol{o}^{(t)}}{\partial \boldsymbol{b}_2}\right)^T \nabla_{\boldsymbol{o}^{(t)}}\mathcal{L} = \sum_t \nabla_{\boldsymbol{o}^{(t)}}\mathcal{L}$$

ILLINOIS

BPTT: Vanishing gradient problem

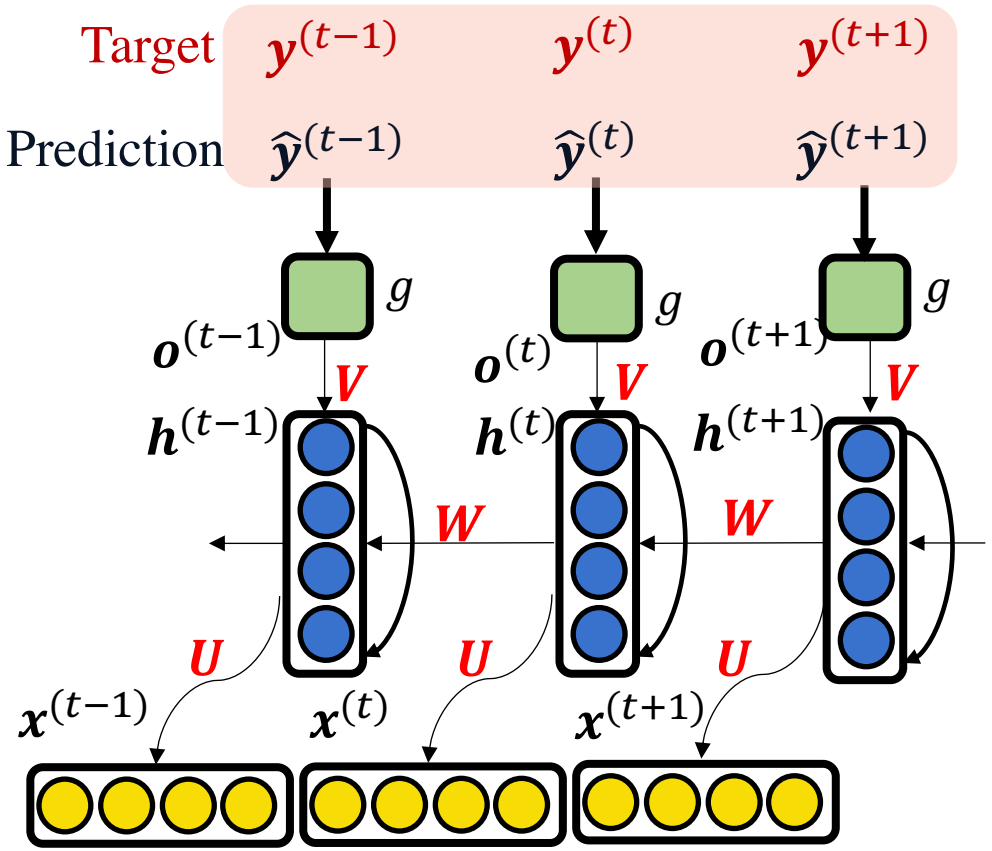

- Gradient can become very small over a long sequence
- Standard RNN will have difficulty to "remember" state from early part of the input sequence

ILLINOIS

Agenda

- RNN Basics
- Learning RNN with Backpropagation Through Time (BPTT)
- Long-Short Term Memory Networks (LSTM)
- Gated Recurrent Unit (GRU)
- Bidirectional RNN
- Sequence-to-Sequence RNN
- Healthcare Applications

Standard RNN

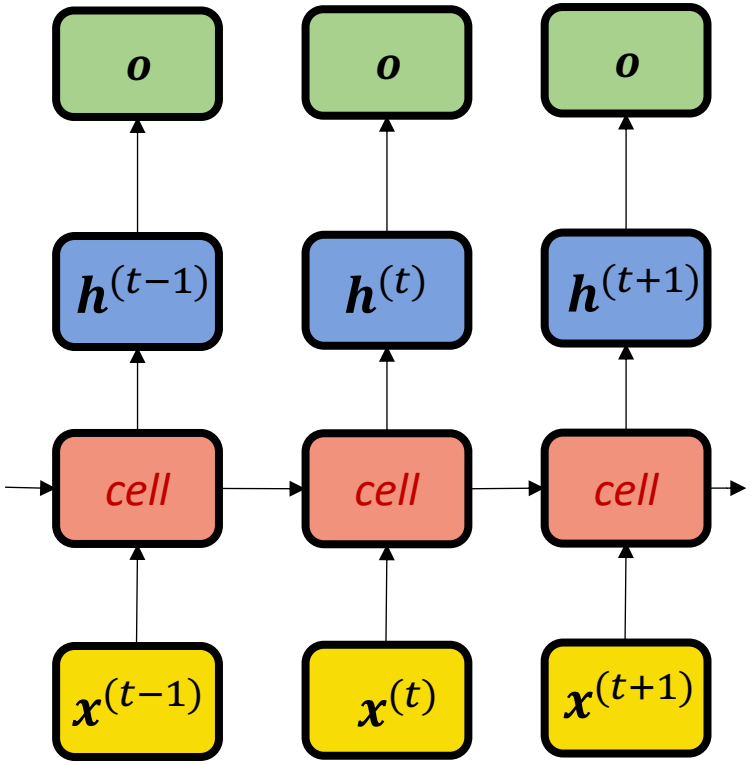

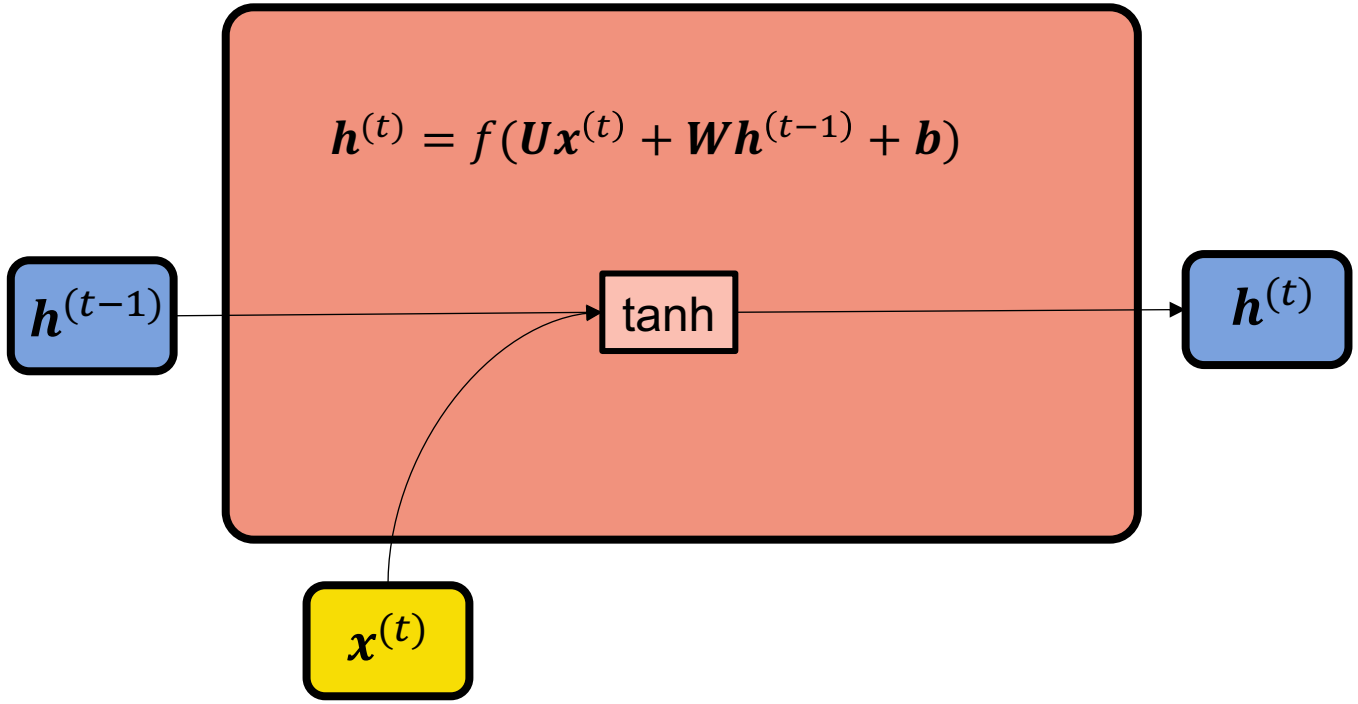

- Standard RNN has a simple computation *cell* from input $\boldsymbol{x}$ to latent state $\boldsymbol{h}$

ILLINOIS

Long short term memory networks (LSTM)

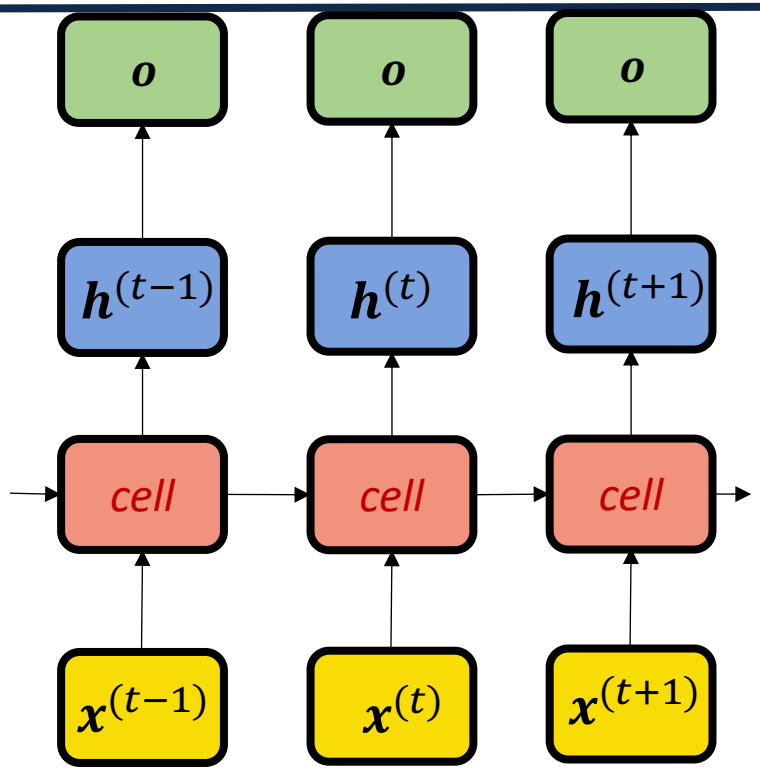

$$h^{(t)} = f(Ux^{(t)} + Wh^{(t-1)} + b)$$

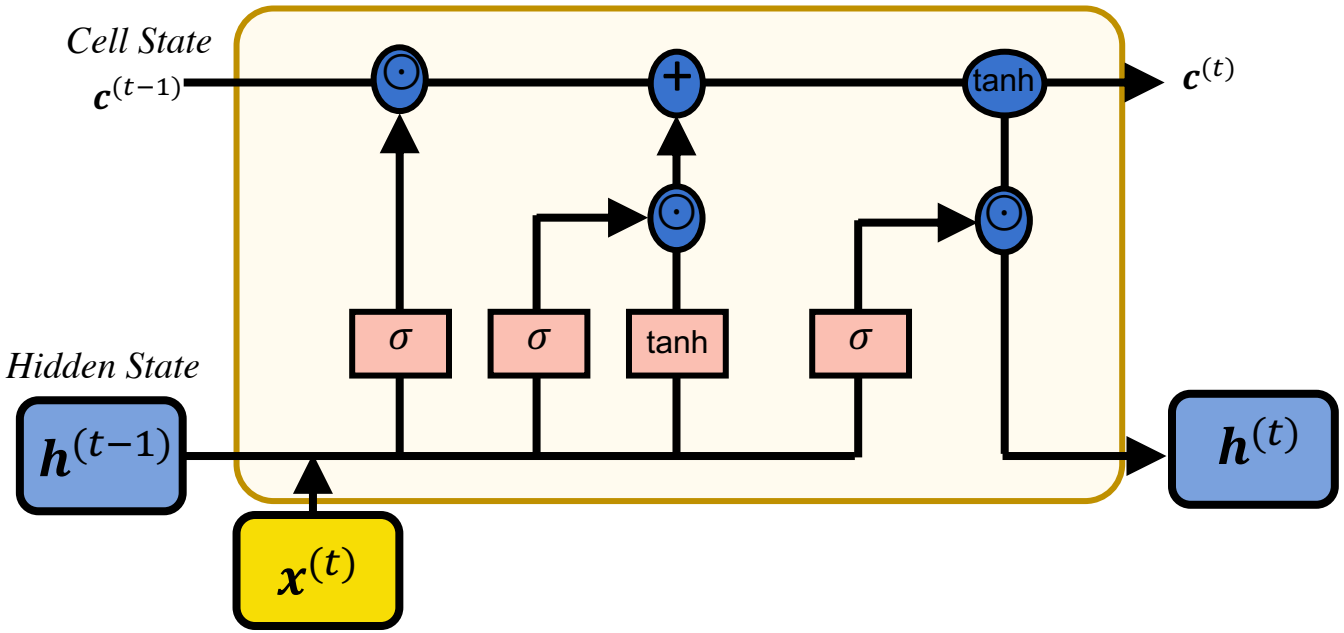

- Standard RNN has a simple computation *cell* from input x to latent state h
- LSTM provides a more sophisticated *cell*

ILLINOIS

LSTM: Cell Structure

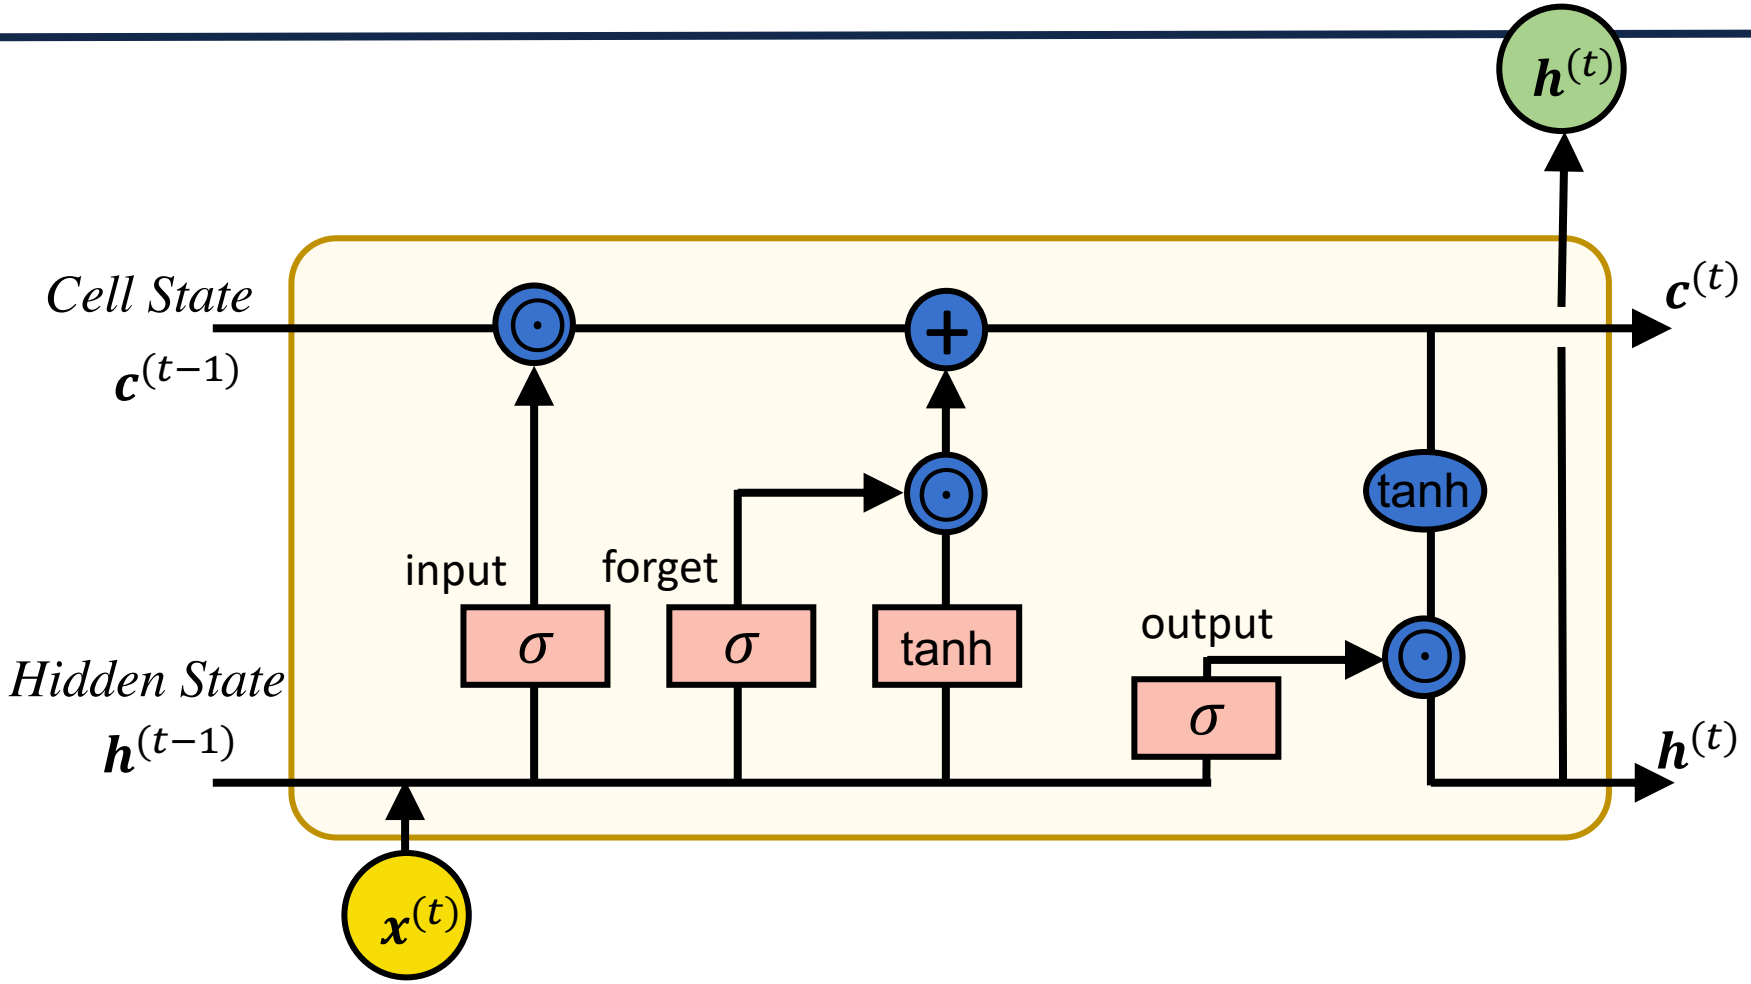

ILLINOIS

LSTM: Cell state

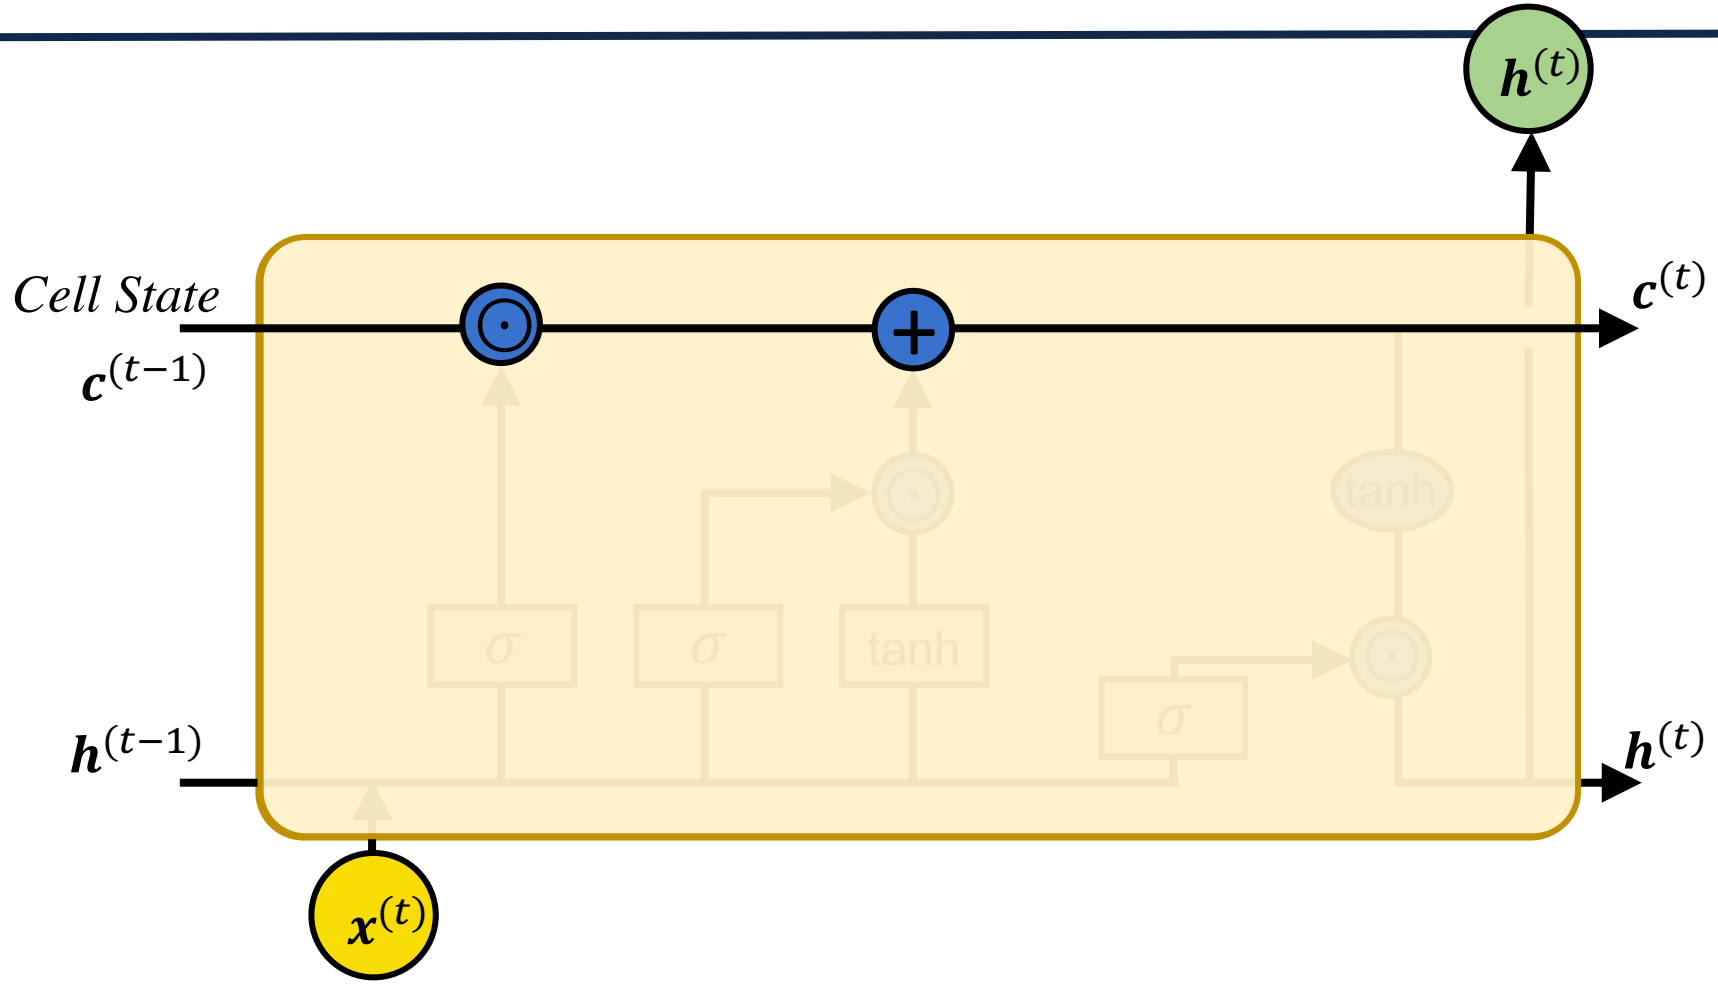

LSTM: Forget Gate

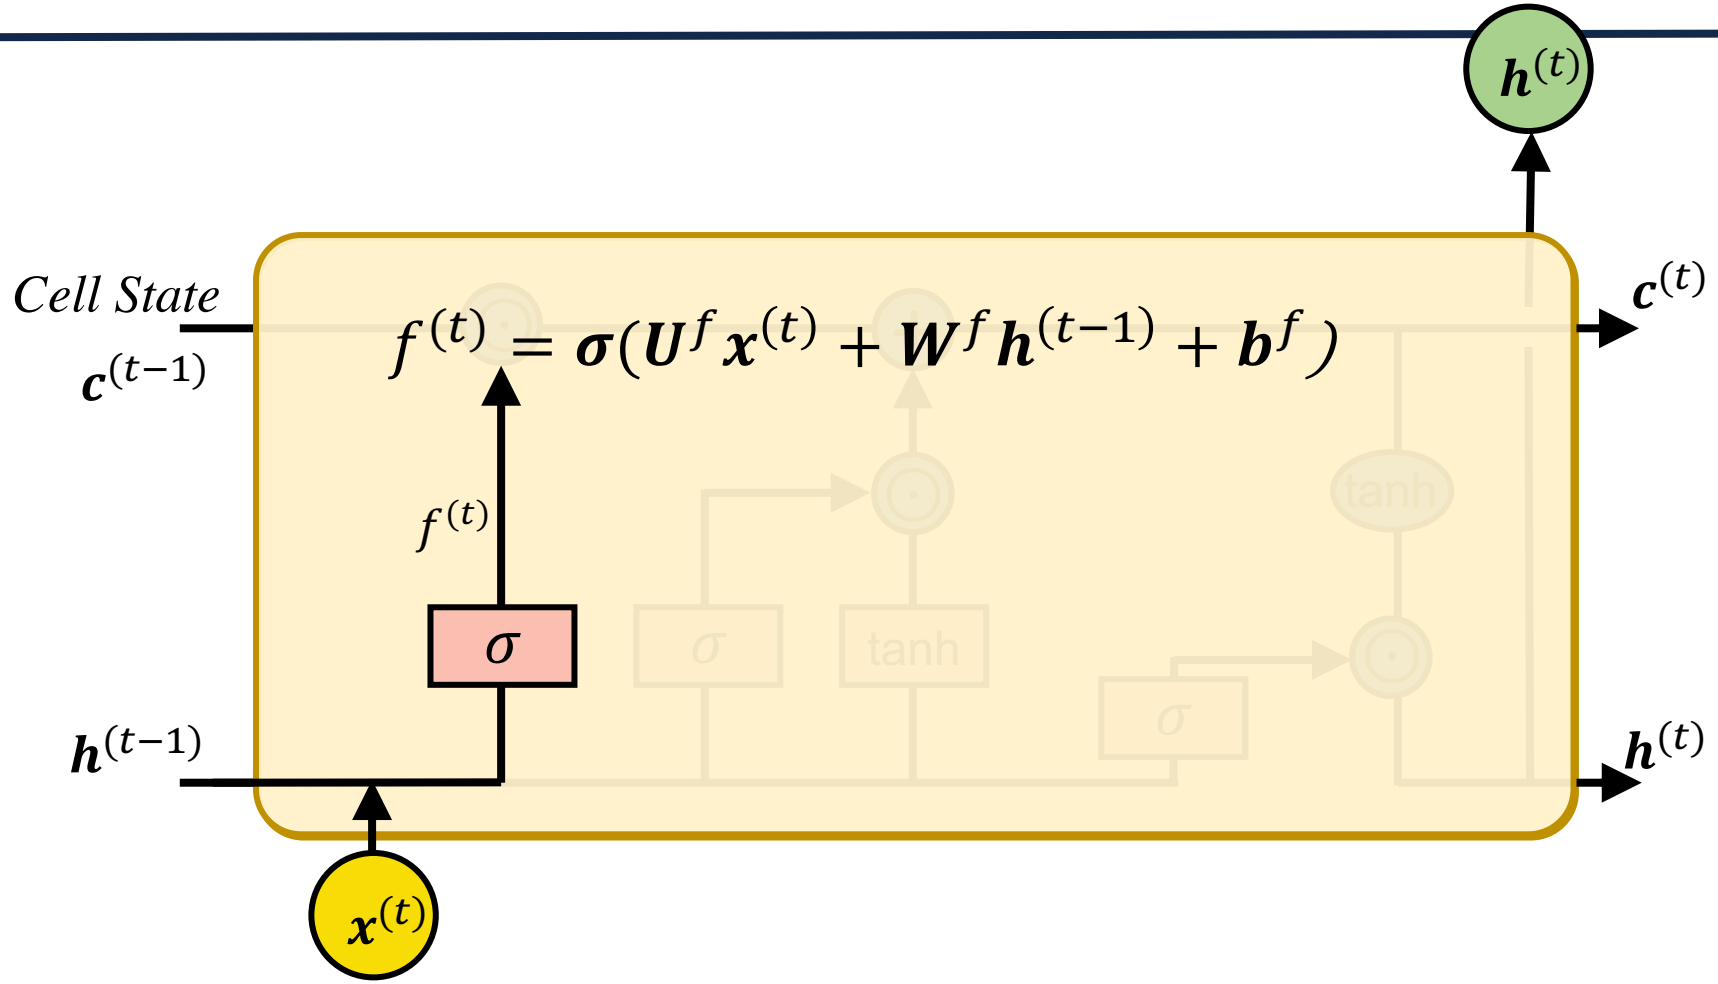

ILLINOIS

LSTM: Input Gate

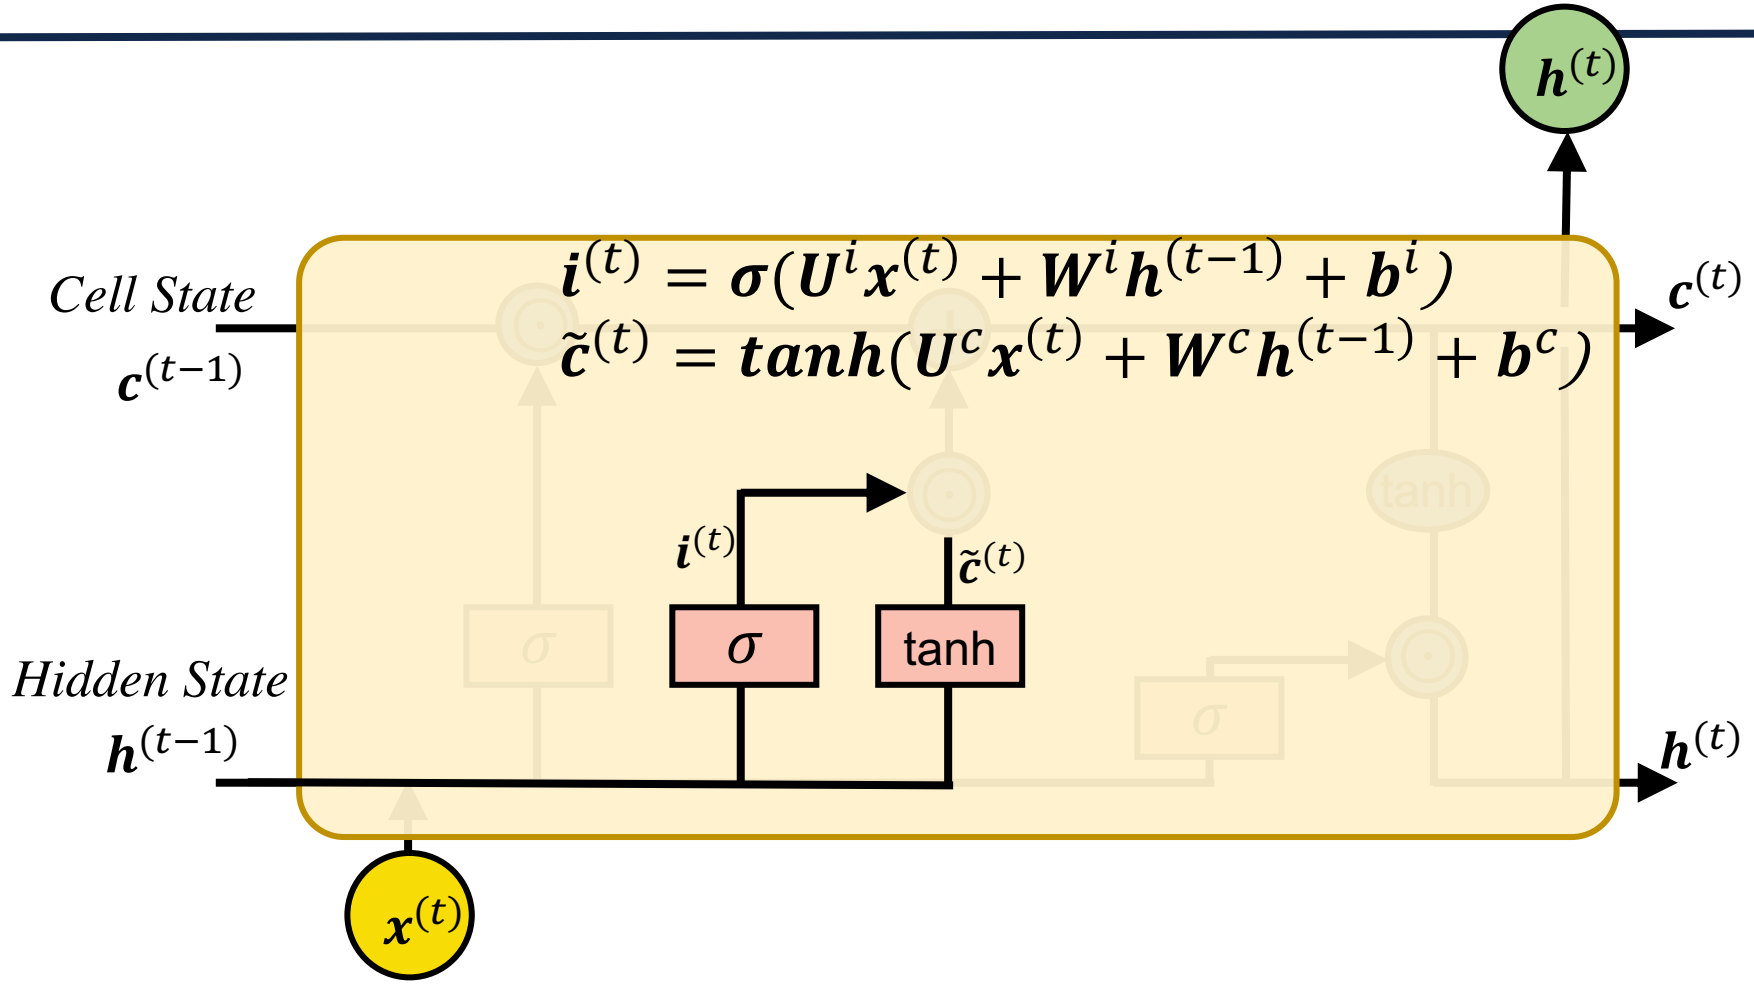

LSTM: Update cell state

$$c^{(t)} = f^{(t)} \odot c^{(t-1)} + i^{(t)} \odot \tilde{c}^{(t)}$$

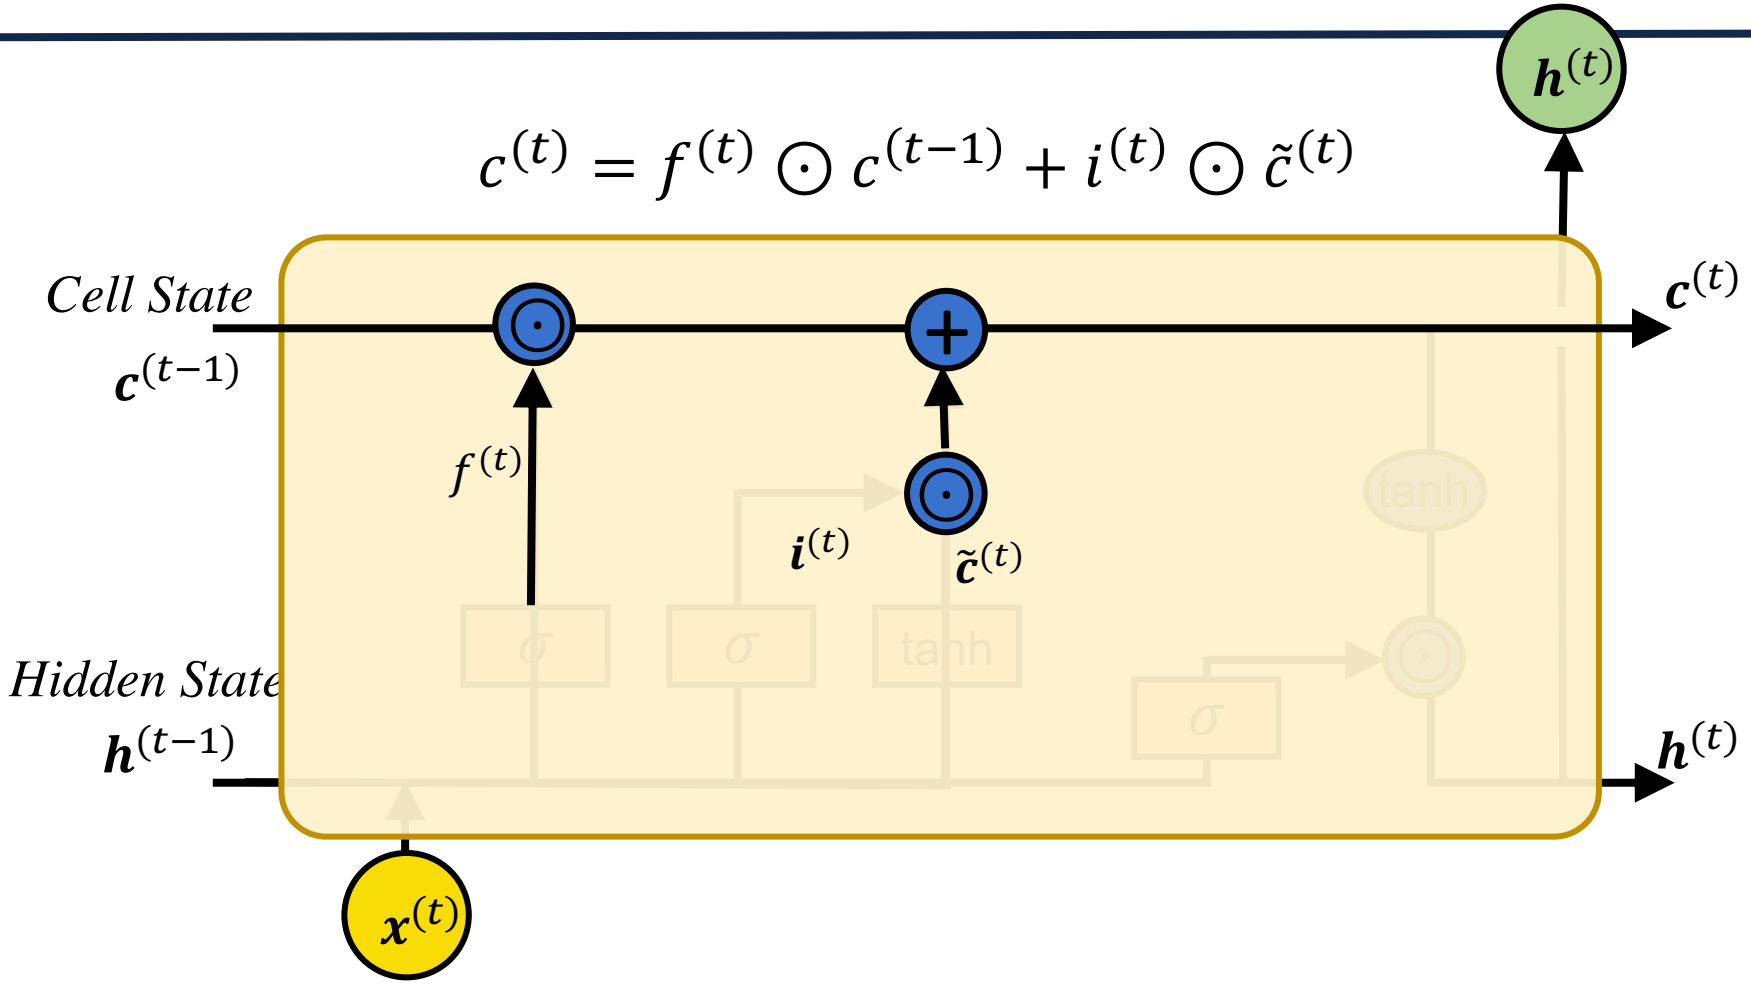

ILLINOIS

LSTM: Output Gate

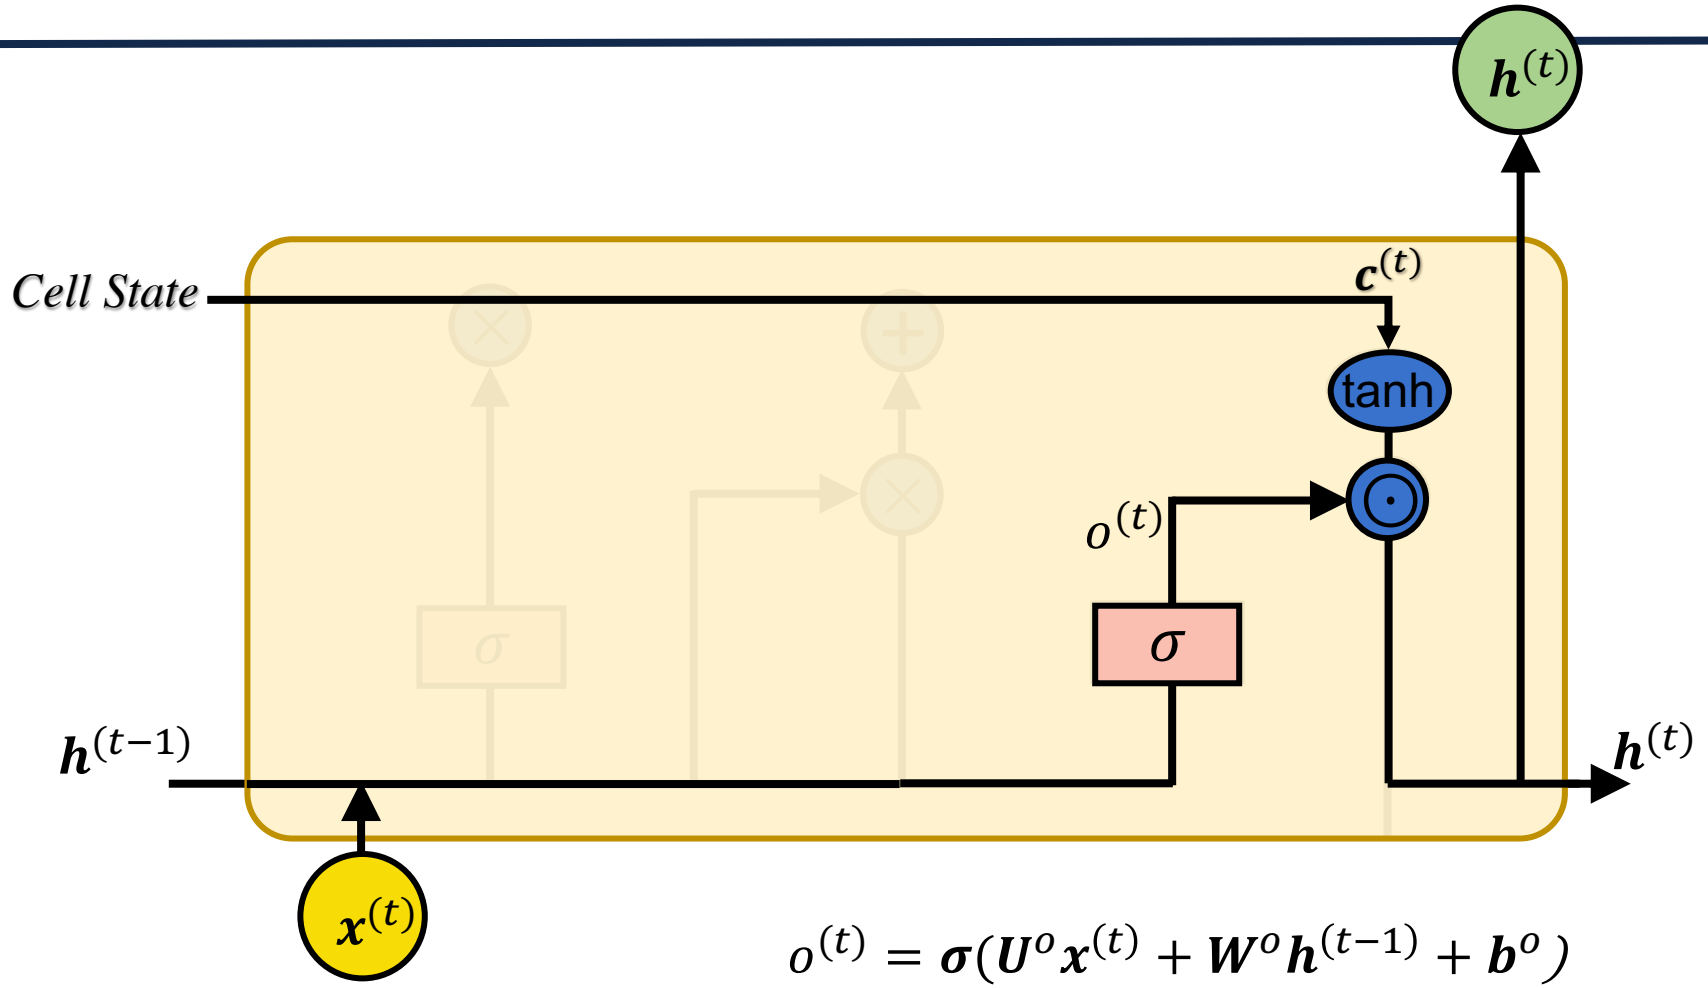

$$o^{(t)} = \boldsymbol{\sigma}(\boldsymbol{U}^o \boldsymbol{x}^{(t)} + \boldsymbol{W}^o \boldsymbol{h}^{(t-1)} + \boldsymbol{b}^o)$$

$$\boldsymbol{h}^{(t)} = o^{(t)} \odot \tanh(c^{(t)})$$

ILLINOIS

LSTM: Cell Structure

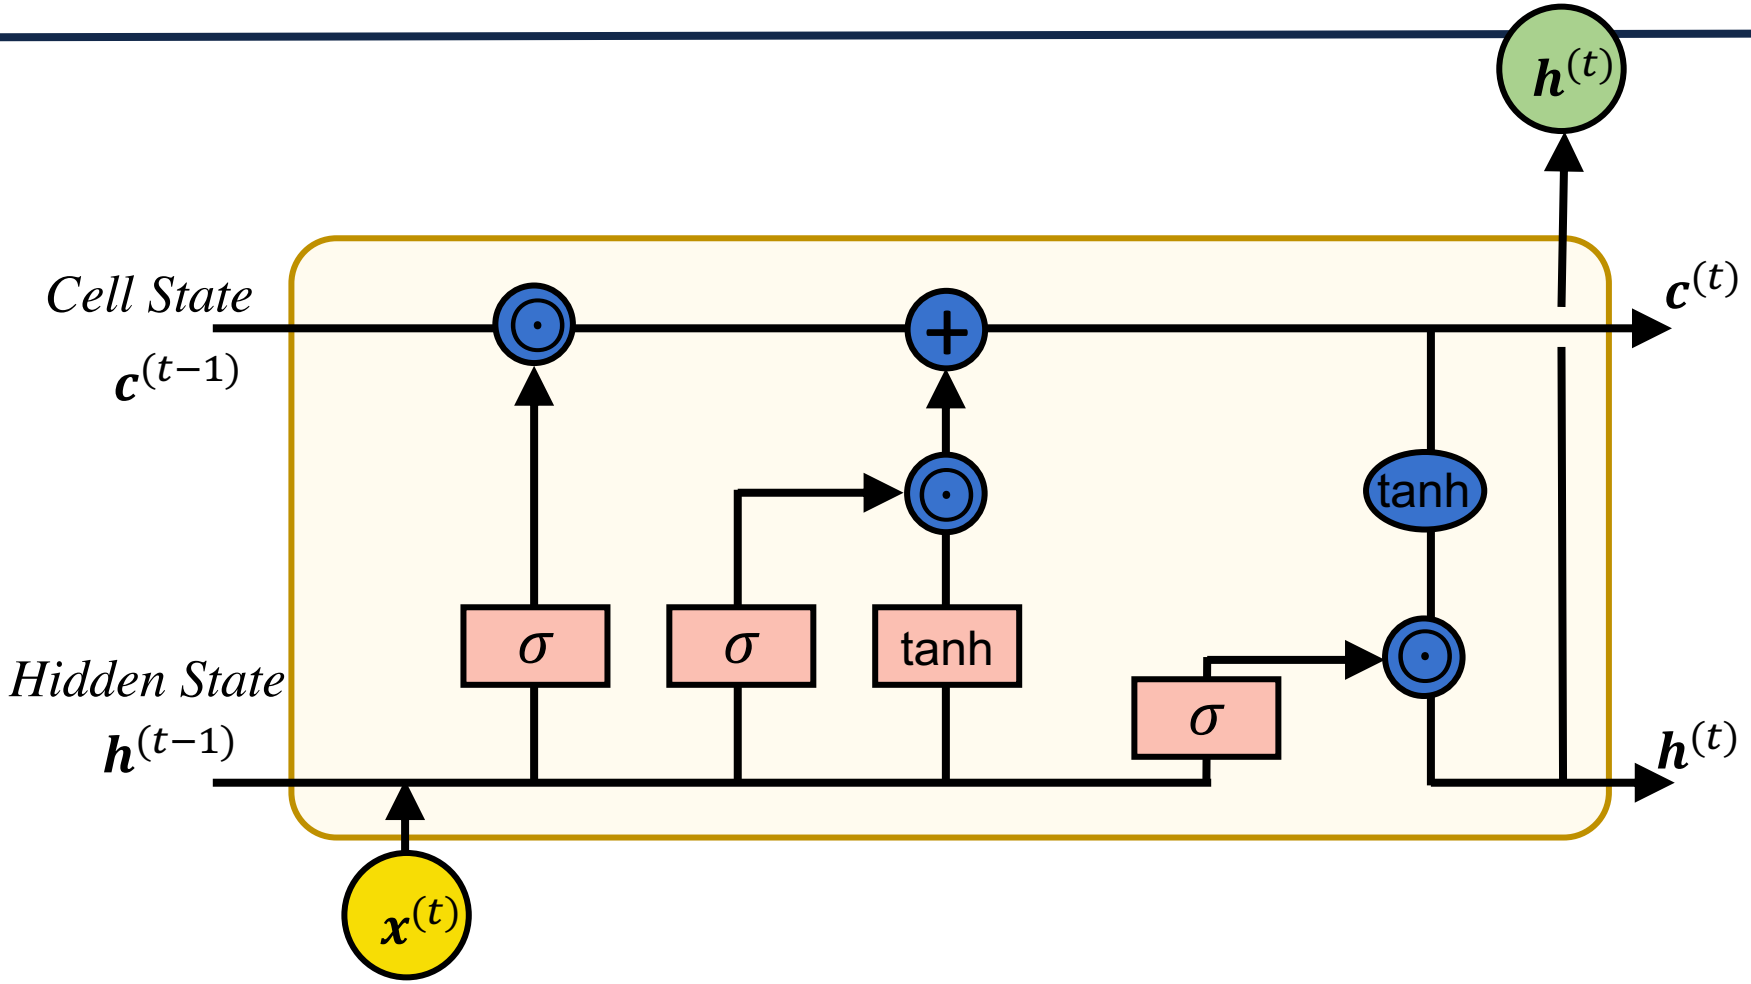

ILLINOIS

Agenda

- RNN Basics
- Learning RNN with Backpropagation Through Time (BPTT)
- Long-Short Term Memory Networks (LSTM)
- Gated Recurrent Unit (GRU)
- Bidirectional RNN
- Sequence-to-Sequence RNN
- Healthcare Applications

ILLINOIS

GRU : Gated Recurrent Unit

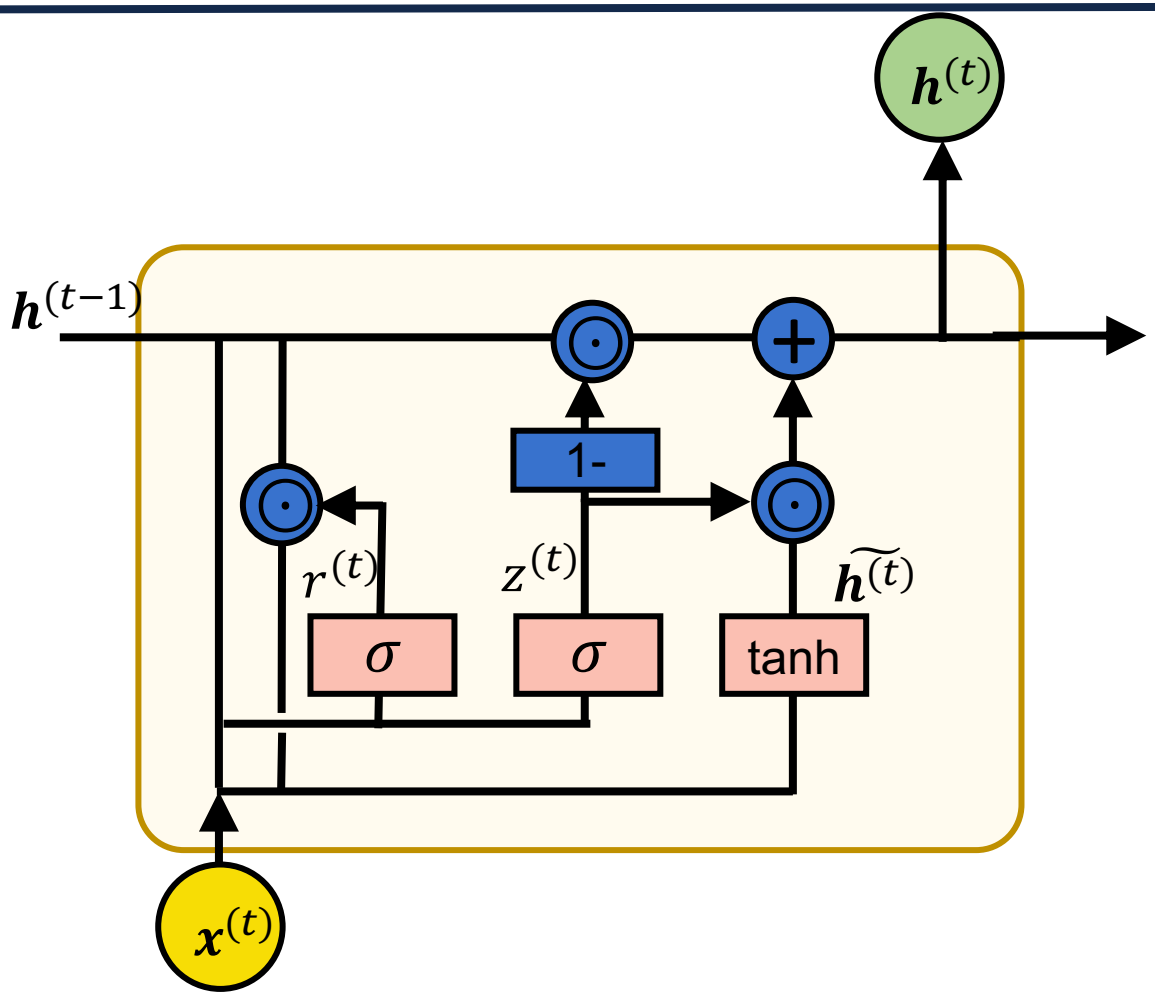

ILLINOIS

GRU: Update Gate

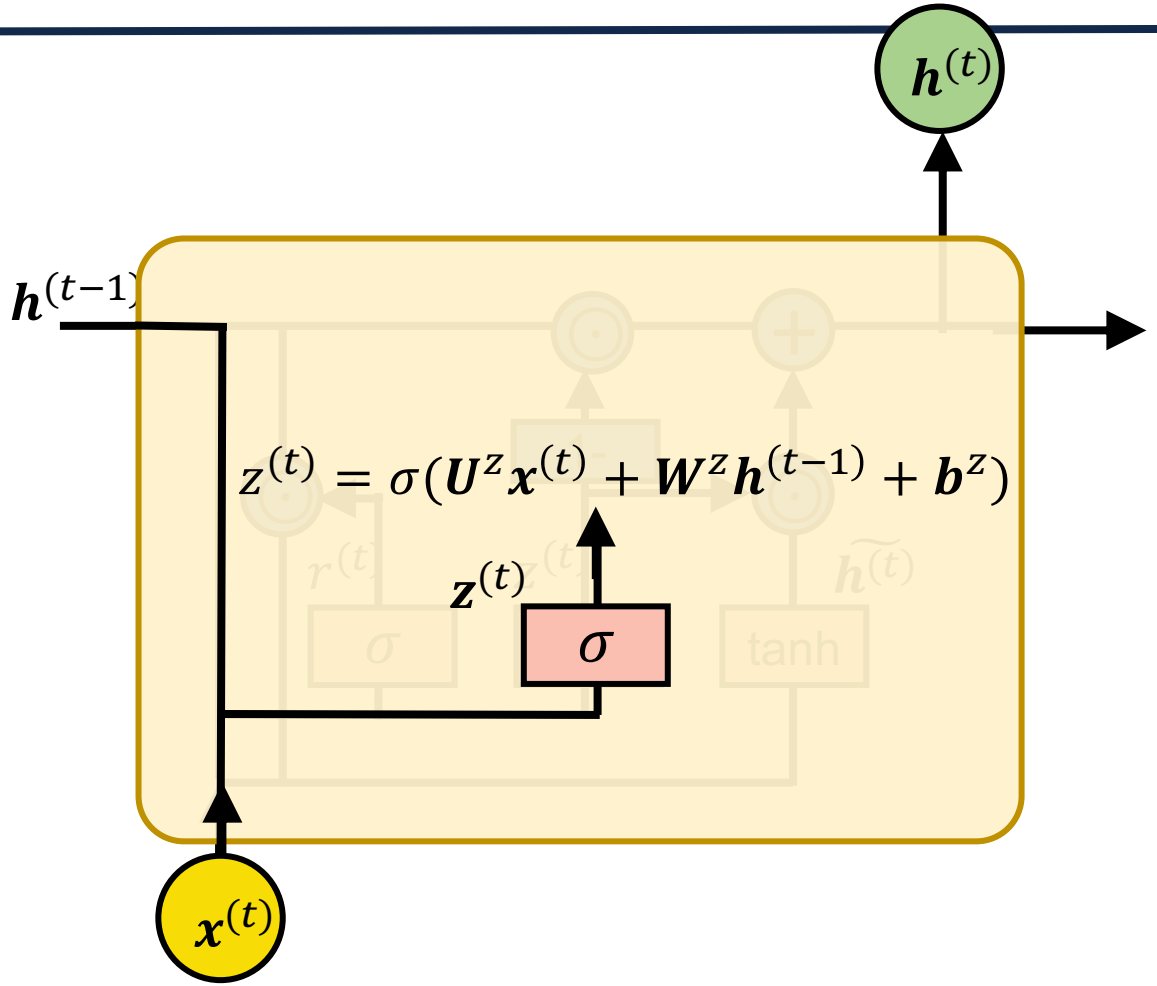

GRU: Reset Gate

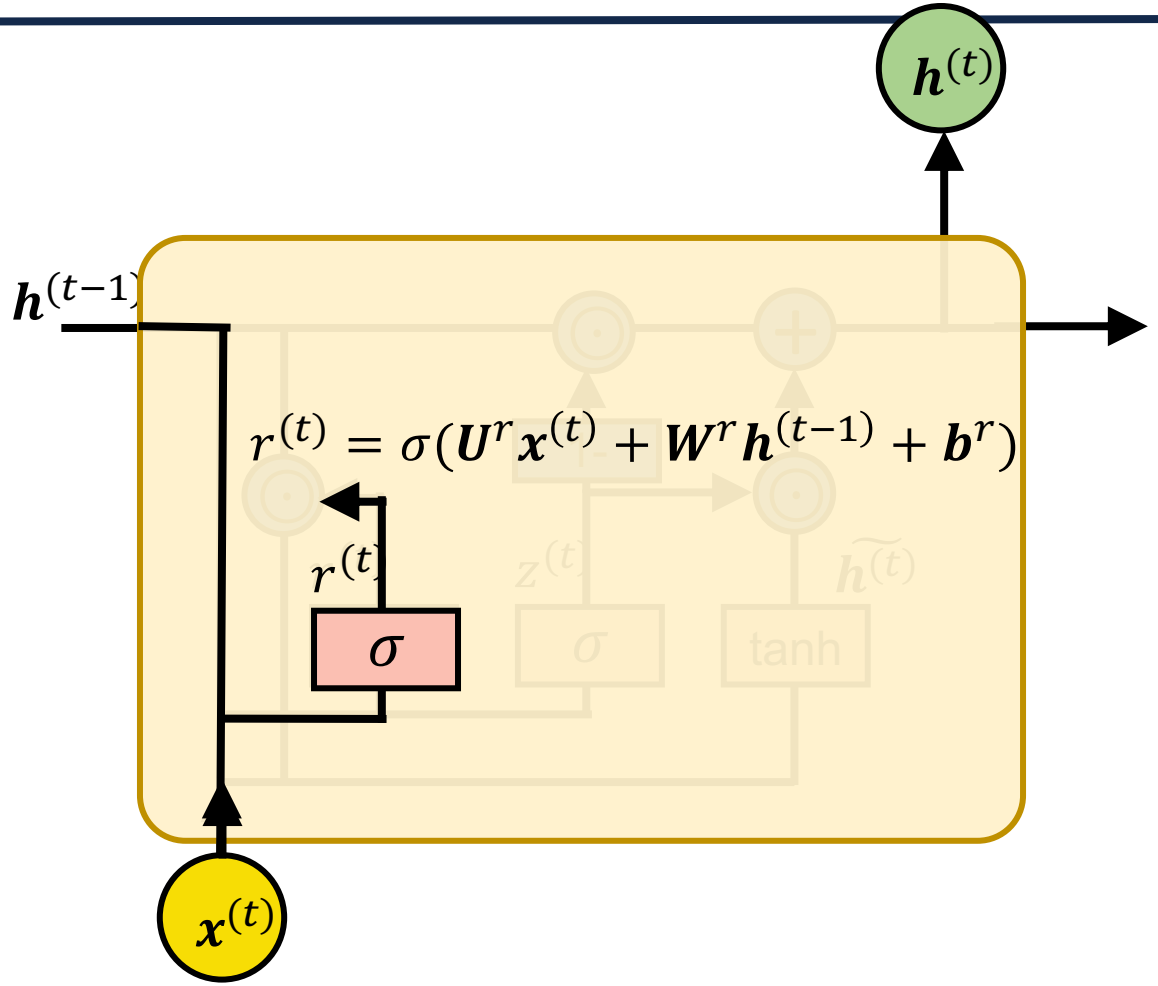

GRU: New information to the hidden state

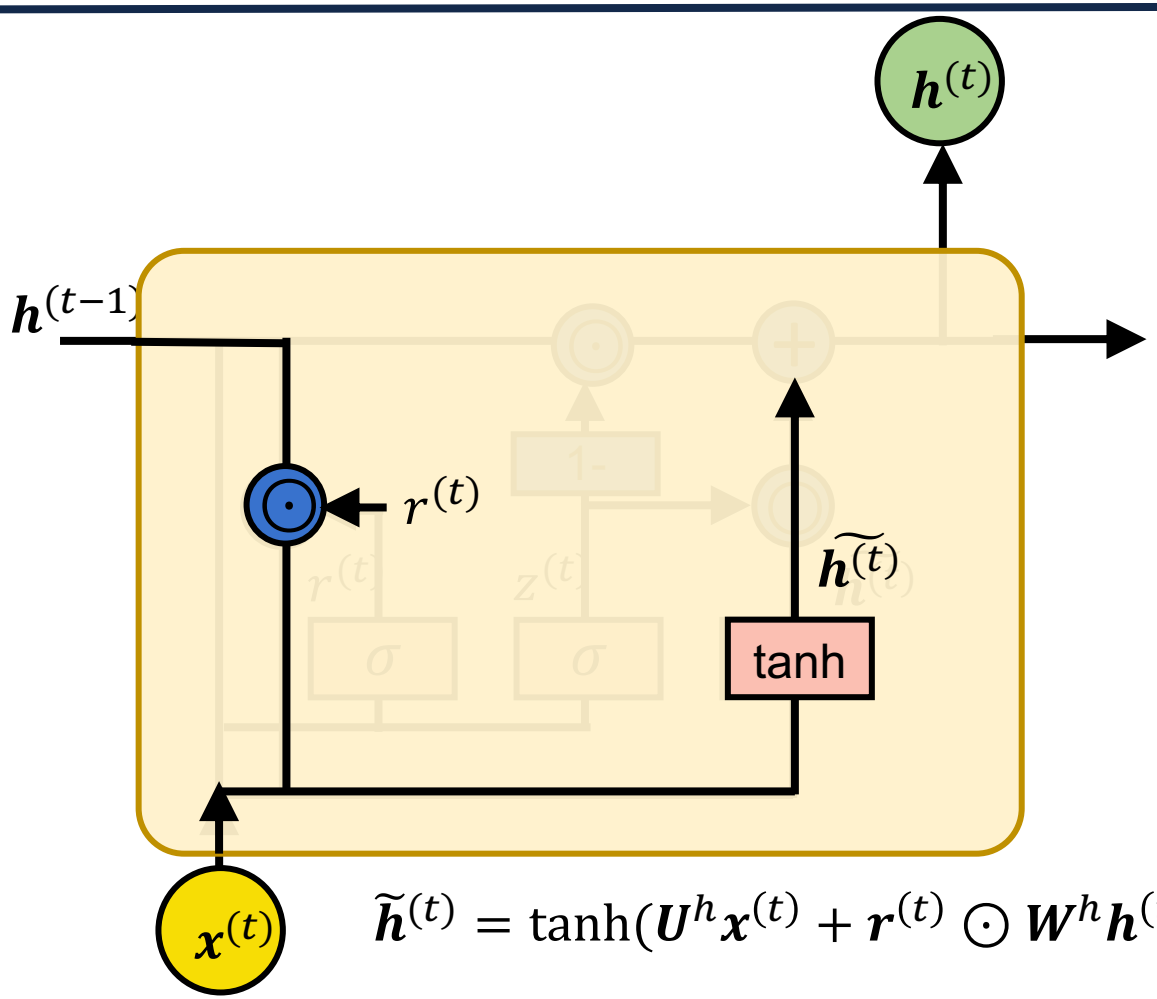

$$\widetilde{\boldsymbol{h}}^{(t)} = \tanh(\boldsymbol{U}^h \boldsymbol{x}^{(t)} + \boldsymbol{r}^{(t)} \odot \boldsymbol{W}^h \boldsymbol{h}^{(t-1)} + \boldsymbol{b}^h)$$

ILLINOIS

GRU: Final New Hidden State

$$\boldsymbol{h}^{(t)} = \left(1 - \boldsymbol{z}^{(t)}\right) \odot \boldsymbol{h}^{(t-1)} + \boldsymbol{z}^{(t)} \odot \widetilde{\boldsymbol{h}}^{(t)}$$

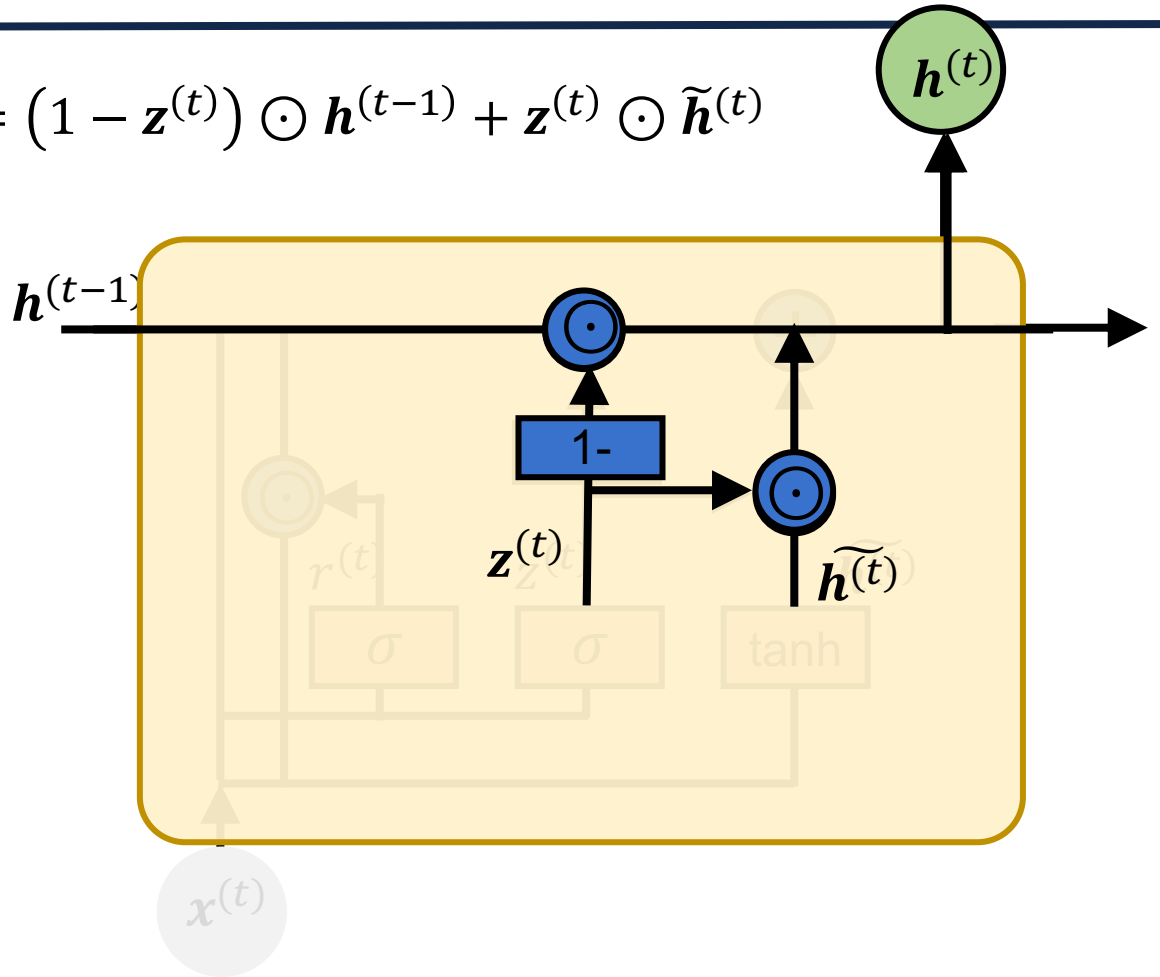

GRU: Summary

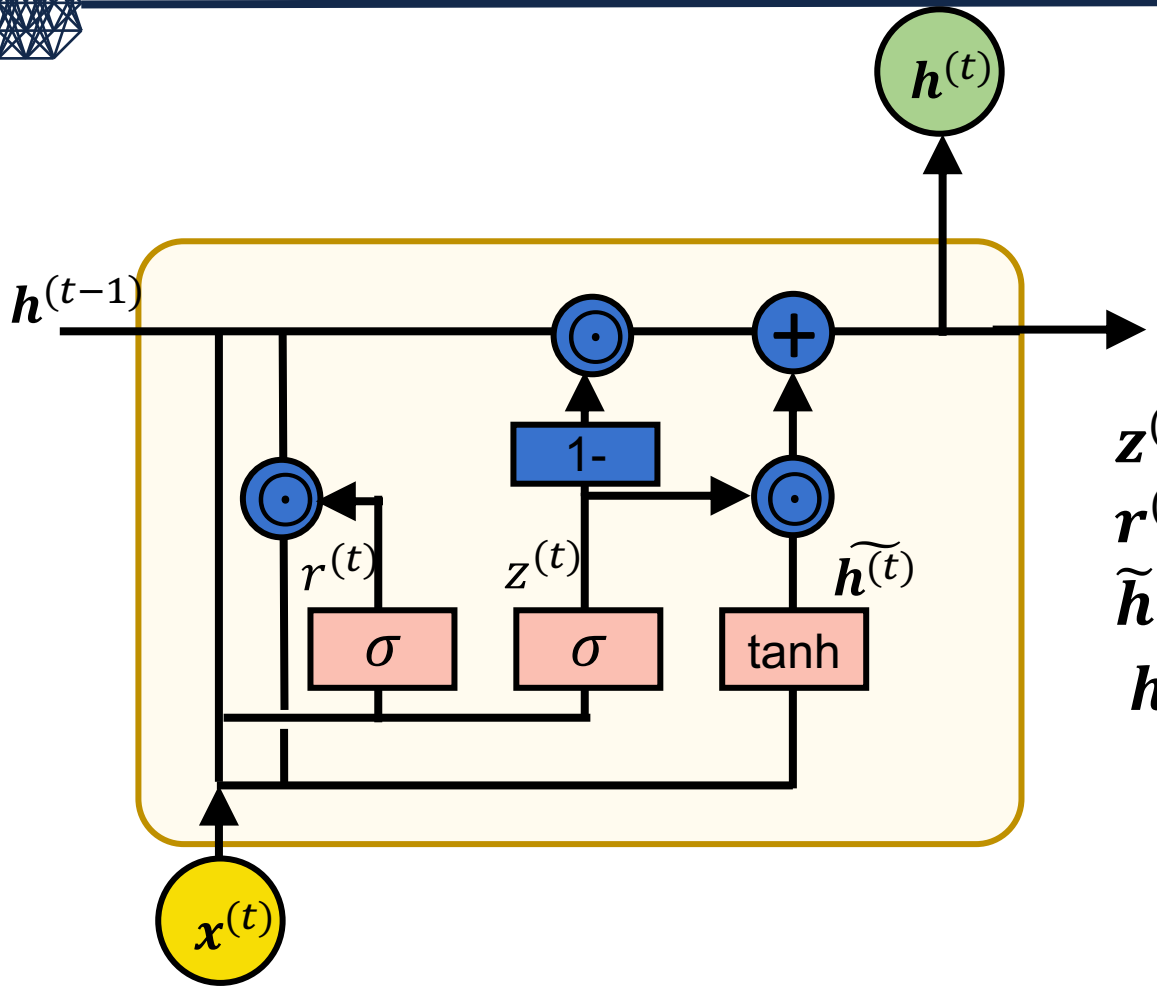

$$z^{(t)} = \sigma(U^z x^{(t)} + W^z h^{(t-1)} + b^z)$$
$$r^{(t)} = \sigma(U^r x^{(t)} + W^r h^{(t-1)} + b^r)$$
$$\tilde{h}^{(t)} = \tanh(U^h x^{(t)} + r^{(t)} \odot W^h h^{(t-1)} + b^h)$$
$$h^{(t)} = (1 - z^{(t)}) \odot h^{(t-1)} + z^{(t)} \odot \tilde{h}^{(t)}$$

ILLINOIS

Agenda

- RNN Basics
- Learning RNN with Backpropagation Through Time (BPTT)
- Long-Short Term Memory Networks (LSTM)
- Gated Recurrent Unit (GRU)
- Bidirectional RNN
- Sequence-to-Sequence RNN
- Healthcare Applications

ILLINOIS

Bidirectional RNN

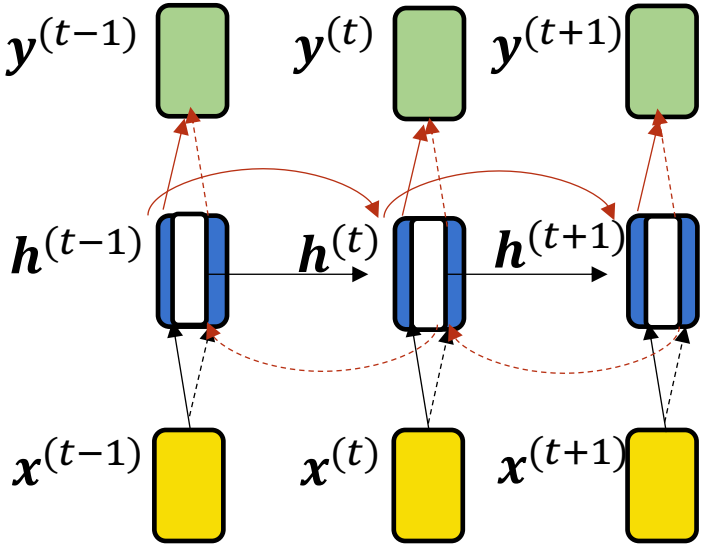

$$\overrightarrow{\boldsymbol{h}}^t = f(\boldsymbol{U}\boldsymbol{x}^{(t)} + \boldsymbol{W}\boldsymbol{h}^{(t-1)} + \boldsymbol{b}_1)$$
$$\overleftarrow{\boldsymbol{h}}^t = f(\boldsymbol{U}\boldsymbol{x}^{(t)} + \boldsymbol{W}\boldsymbol{h}^{(t+1)} + \boldsymbol{b}_1)$$
$$y^{(t)} = g(\boldsymbol{V}\left[\overrightarrow{\boldsymbol{h}}^t; \overleftarrow{\boldsymbol{h}}^t\right] + \boldsymbol{b}_2)$$

ILLINOIS

Agenda

- RNN Basics
- Learning RNN with Backpropagation Through Time (BPTT)
- Long-Short Term Memory Networks (LSTM)
- Gated Recurrent Unit (GRU)
- Bidirectional RNN
- Sequence-to-Sequence RNN
- Healthcare Applications

ILLINOIS

Encoder-Decoder Sequence-to-Sequence Model

$$P_\theta(\mathcal{Y}|\mathcal{X}) = \prod_{j=1}^{J+1} P_\theta(\boldsymbol{y}_j|\boldsymbol{y}_{j-1}, \cdots, \boldsymbol{y}_1, \mathcal{X})$$

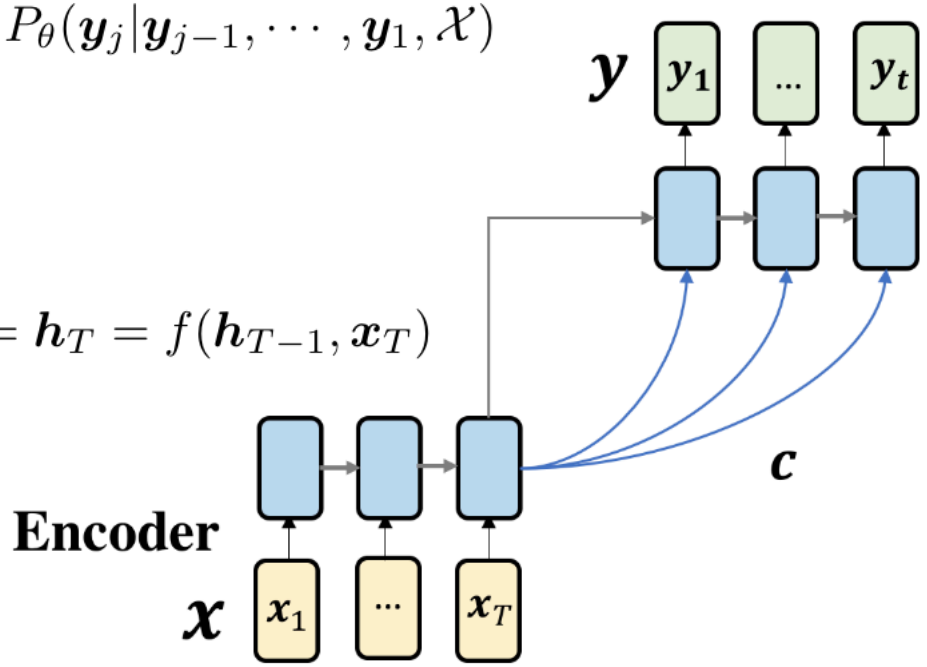

$$P(\boldsymbol{y}_t|\boldsymbol{y}_{t-1}, \cdots, \boldsymbol{y}_1, \boldsymbol{c}) = g(\boldsymbol{h}_t, \boldsymbol{y}_{t-1}, \boldsymbol{c})$$

ILLINOIS

Agenda

- RNN Basics
- Learning RNN with Backpropagation Through Time (BPTT)
- Long-Short Term Memory Networks (LSTM)
- Gated Recurrent Unit (GRU)
- Bidirectional RNN
- Sequence-to-Sequence RNN
- Healthcare Applications

ILLINOIS

Using Recurrent Neural Network Models For Early Detection Of Heart Failure Onset

How to model temporal relations in EHR

Edward Choi, Andy Schuetz, Walter Stewart, Jimeng Sun. Using Recursive Neural Network Models for Early Detection of Heart Failure Onset, JAMIA 2017

Heart Failure Prediction with RNN

- Given a patient record, predict if the patient will be diagnosed with heart failure in the future

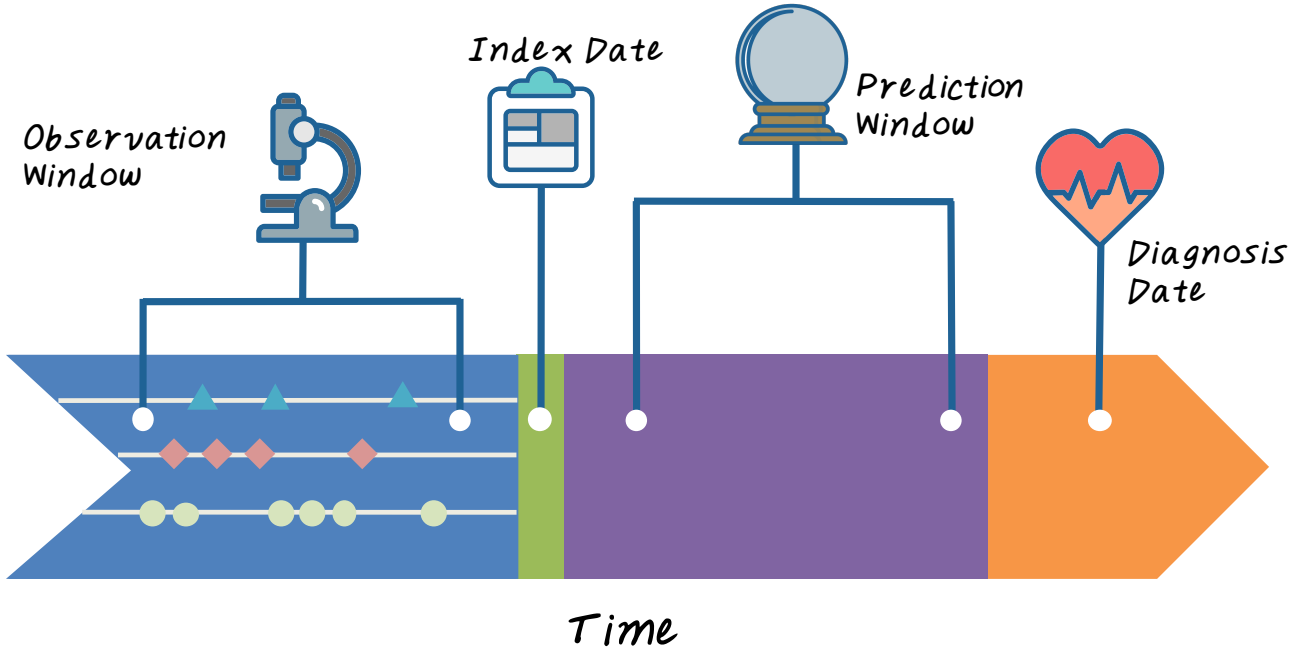

ILLINOIS

Heart Failure Prediction with RNN

- Input sample **x**
  - Patient record over time
  - Diagnosis codes, medication codes, procedure codes

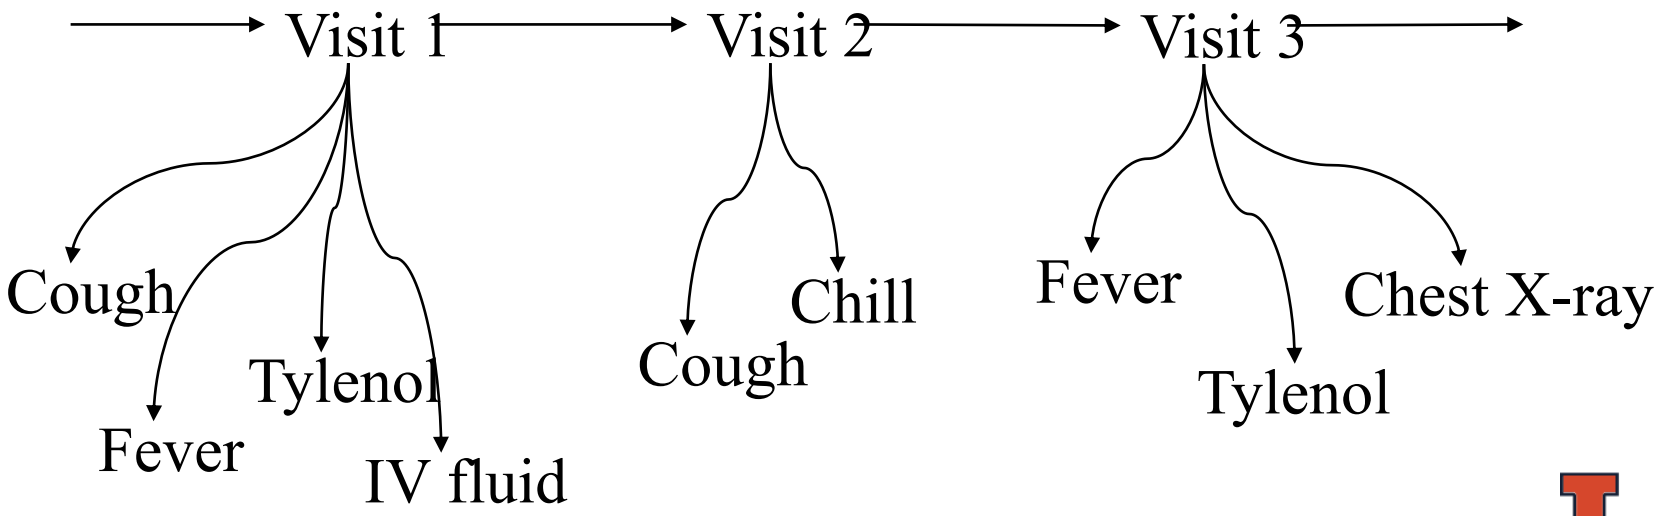

ILLINOIS

Heart Failure Prediction with RNN

- Input sample **x**
  - Patient record over time
  - Diagnosis codes, medication codes, procedure codes

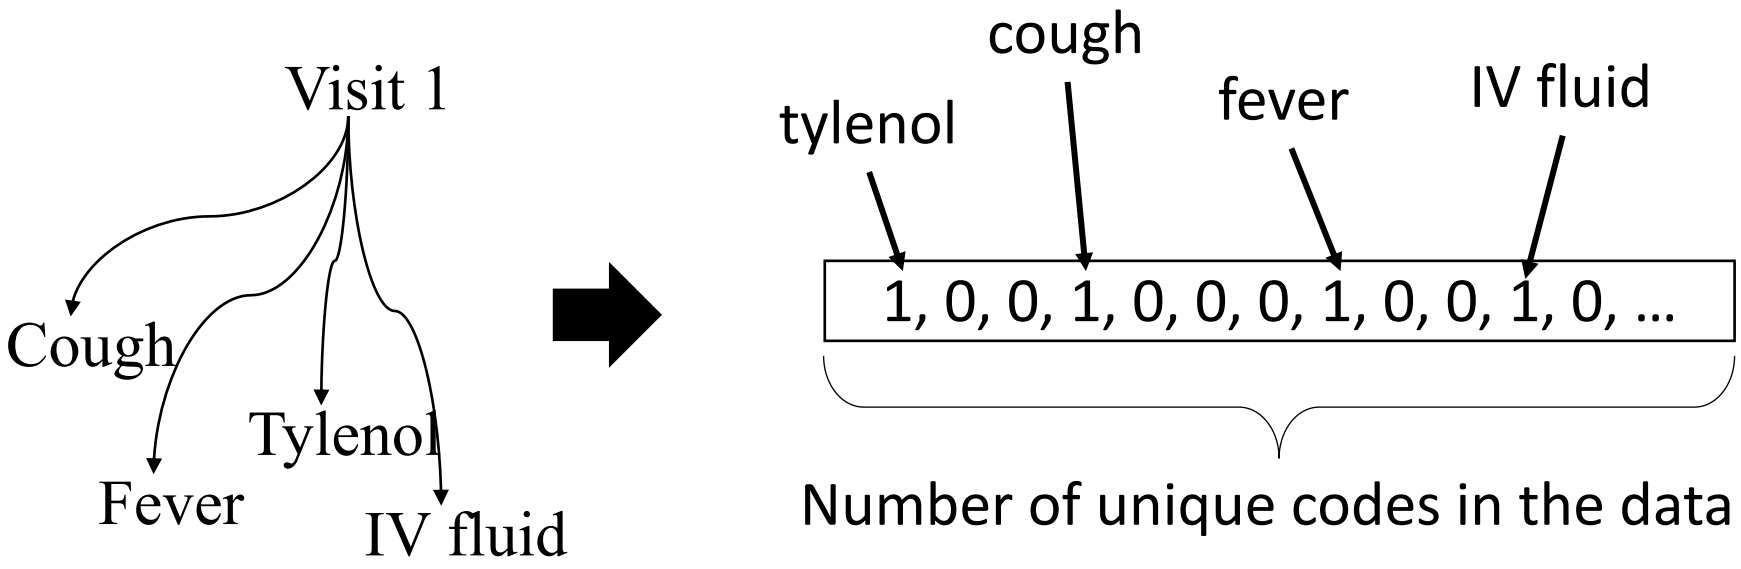

ILLINOIS

Heart Failure Prediction with RNN

- Feed visits into the RNN
  - One visit at each timestep
  - Use the final hidden layer to predict HF

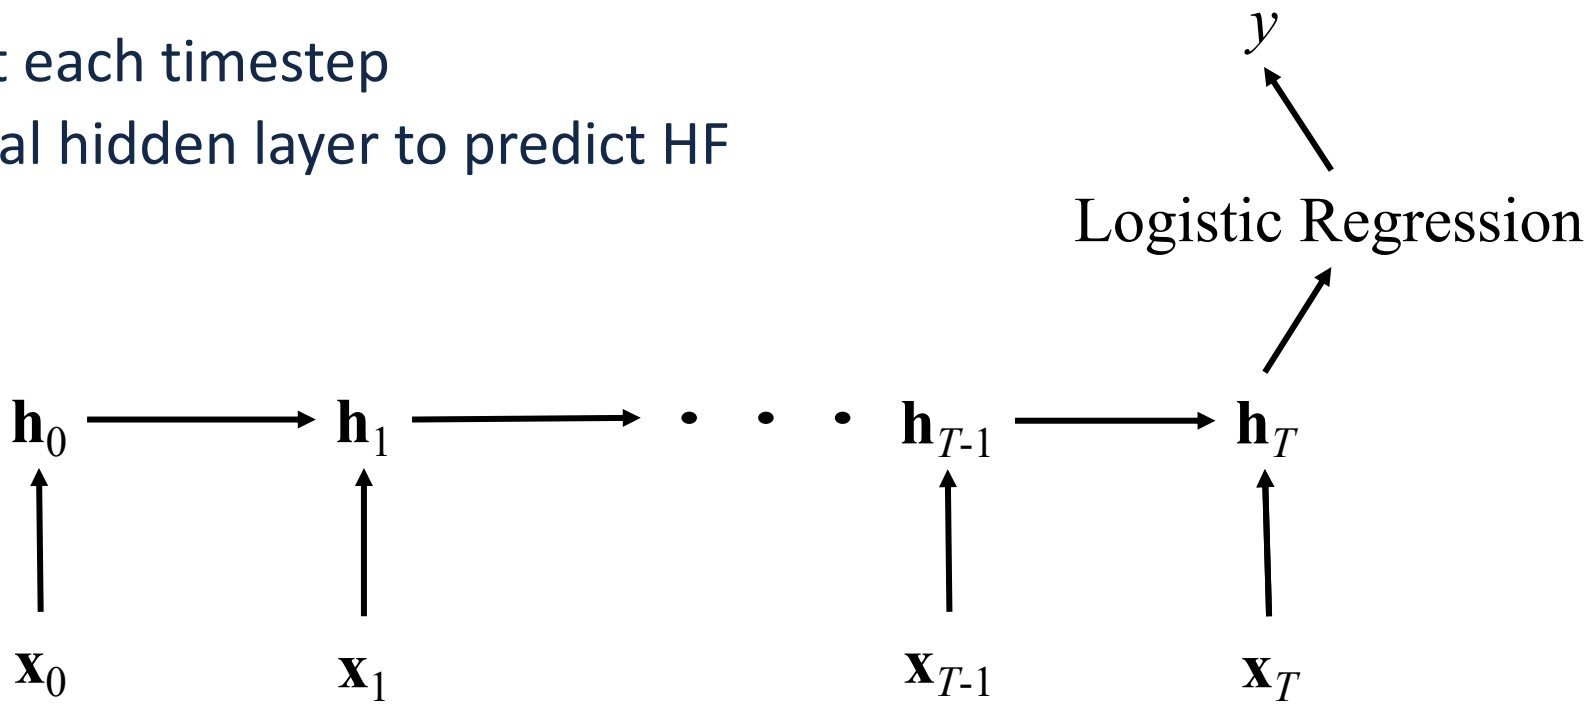

ILLINOIS

Heart Failure Prediction with RNN

Data

34K patients from Sutter Health

4K cases, 30K controls

18-months observation window

Case-control selection criteria

Age (40-85)

Types of diagnoses received

Number of hospital visits

Time span between diagnoses

ILLINOIS

Heart Failure Prediction with RNN

- Prediction performance

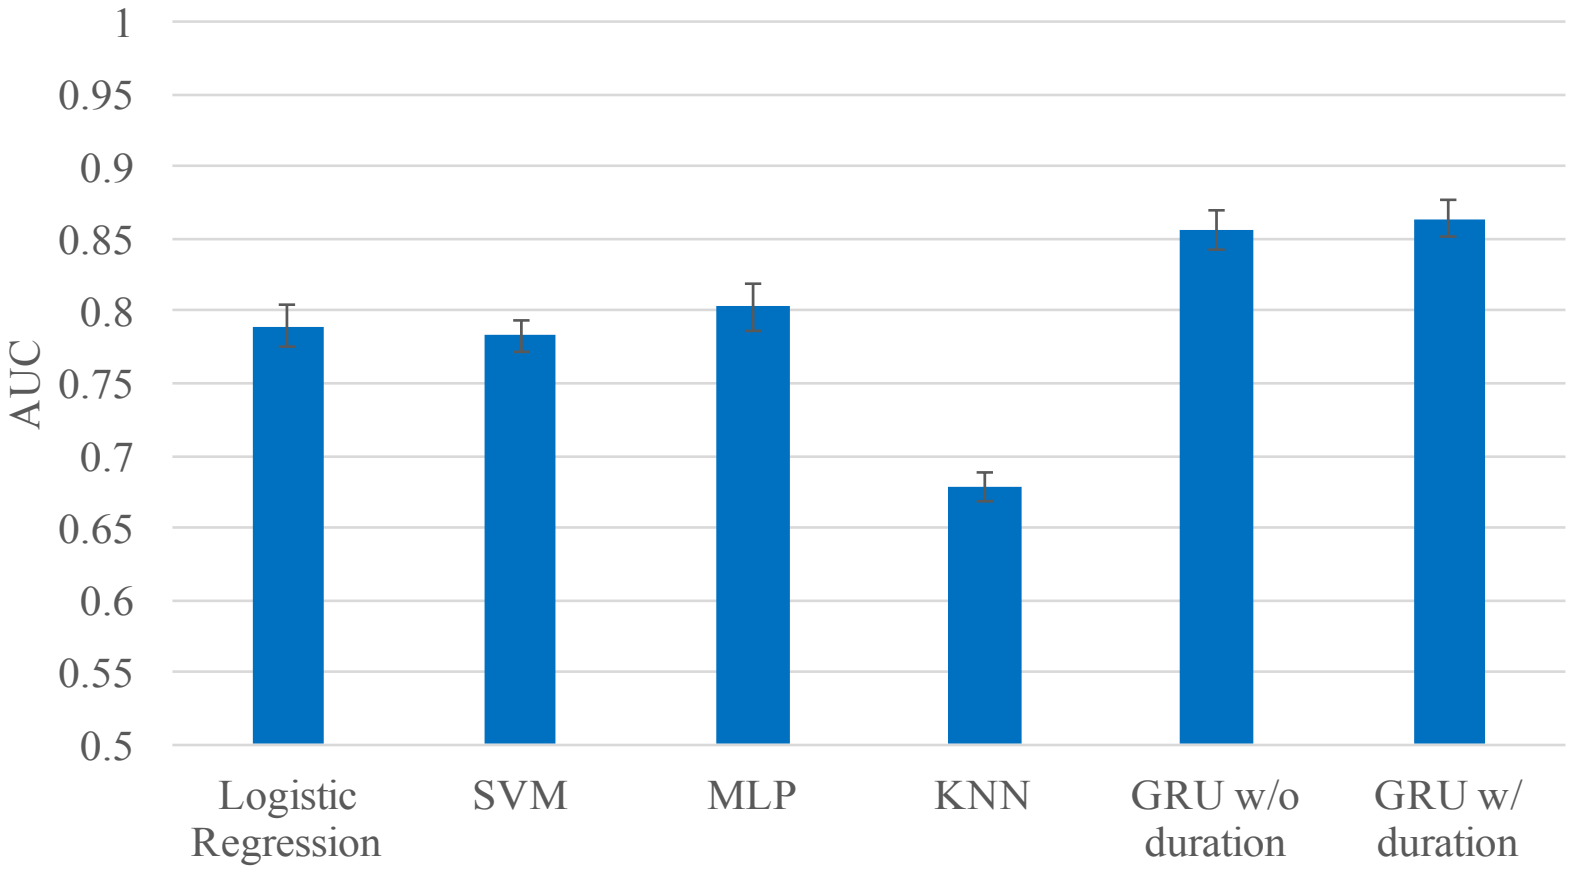

ILLINOIS

Doctor AI: Predicting Clinical Events via Recurrent Neural Networks

Edward Choi, Mohammad Taha Bahadori, Andy Schuetz, Walter F. Stewart, Jimeng Sun

Machine Learning for Healthcare Conference, 2016

ILLINOIS

Doctor AI: Background

- Disease progression modeling

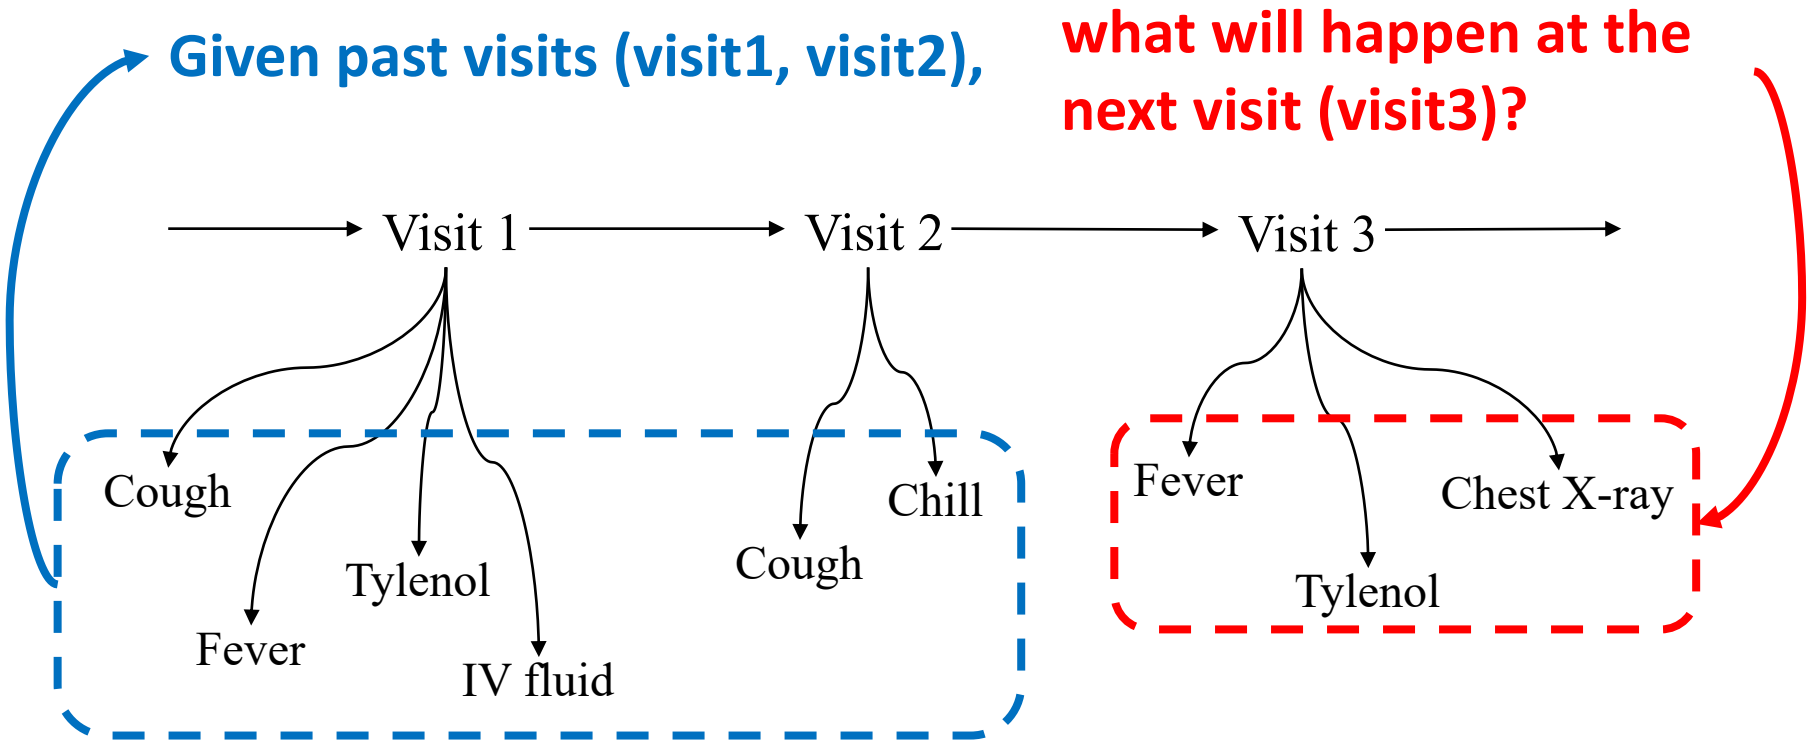

ILLINOIS

Doctor AI: Model

- Feed visits into the RNN
  - One visit at each timestep.
  - Predict next events at each timestep.

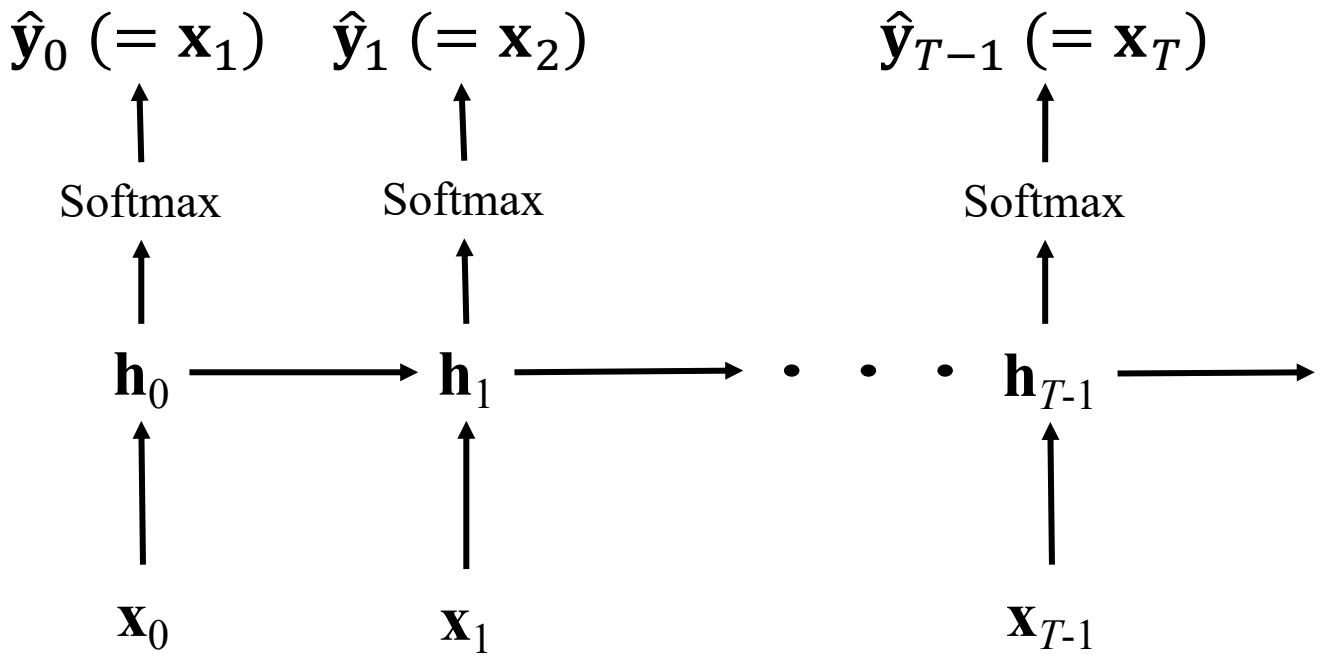

ILLINOIS

Doctor AI: Data

- 260K patients from Sutter Health
- Patient records over 10 years
- Input codes
  - Diagnosis codes, medication codes, procedure codes (38,000 codes)
- Output labels
  - 1,183 diagnosis codes

ILLINOIS

Doctor AI: Sequential Prediction

- Predicting diagnoses in the next visit

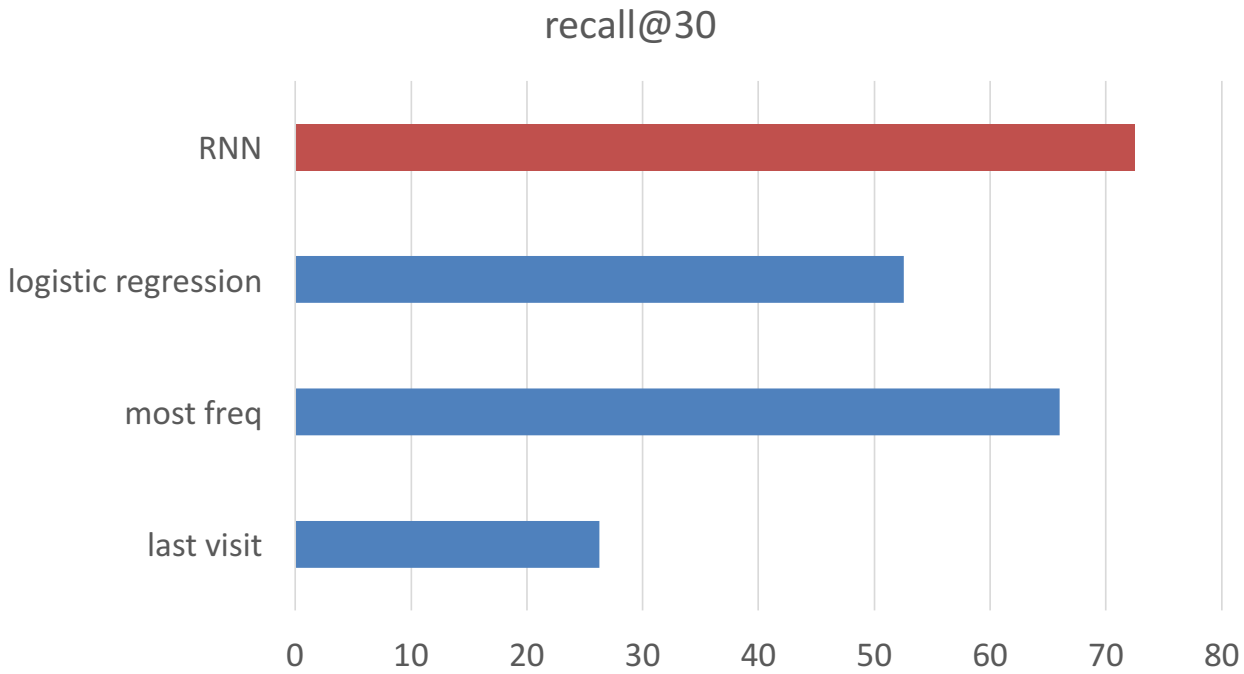

$$\text{top-}k \text{ recall} = \frac{\#\text{ of true positives in the top } k \text{ predictions}}{\#\text{ of true positives}}$$

ILLINOIS

Doctor AI: Knowledge Transfer

- Generalize RNN model from one hospital to another

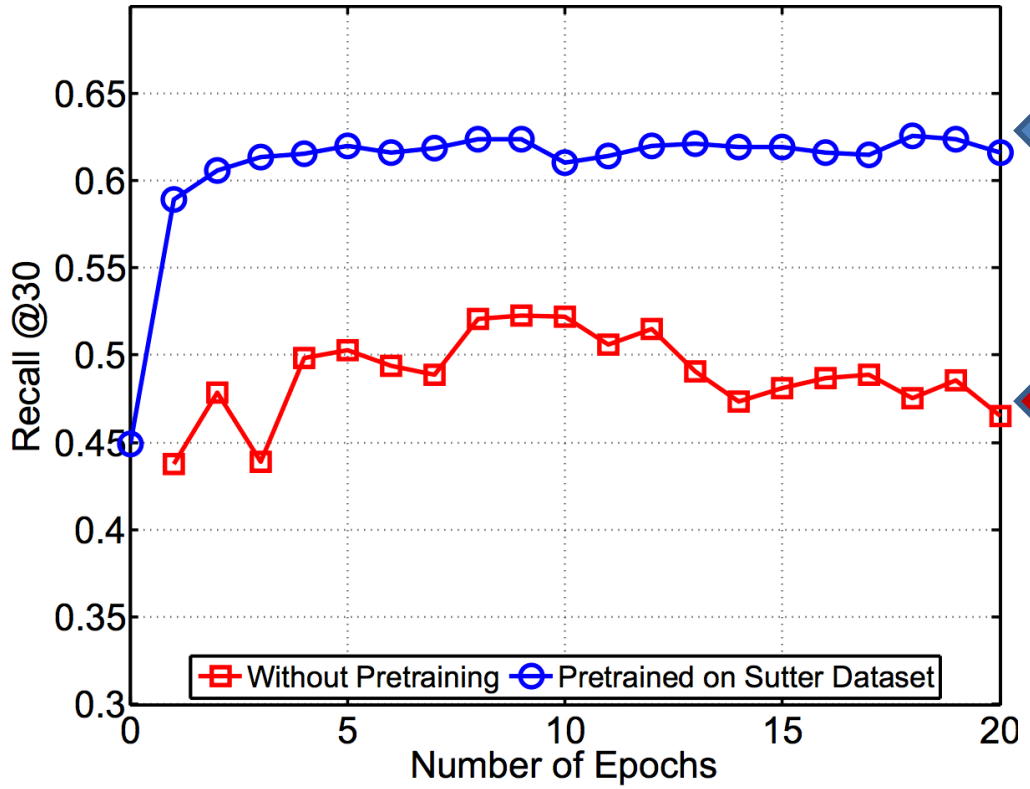

Perform better when warm start the existing model

Perform worse when cold-start from random initialization

ILLINOIS